

Contents

1	The Shear-Vorticity Identity and Spectral Concentration in SQG Front Dynamics	2
1.1	Abstract	2
1.2	1. Introduction	3
1.3	2. The Shear-Vorticity Identity	4
1.3.1	2.1 Statement	4
1.3.2	2.2 Proof	4
1.3.3	2.3 Physical interpretation	5
1.4	3. Consequences for the Curvature Budget	5
1.4.1	3.1 Level-set curvature evolution	5
1.4.2	3.2 Application of the identity	5
1.5	4. Wavevector Rotation and Spectral Concentration	6
1.5.1	4.1 Material gradient dynamics	6
1.5.2	4.2 The identity eliminates the forcing	6
1.5.3	4.3 Self-reinforcing cycle	6
1.6	5. Numerical Verification	7
1.6.1	5.1 Setup	7
1.6.2	5.2 Identity verification	7
1.6.3	5.3 Spectral concentration	7
1.6.4	5.4 Local amplitude of $f = S_{nt} - \omega/2$ near x^*	7
1.6.5	5.5 Two dynamical regimes	8
1.7	6. Implications for SQG Regularity	9
1.7.1	6.1 The selection rule	9
1.7.2	6.2 Self-consistent bound	9
1.7.3	6.3 Cascade argument	9
1.7.4	6.4 The local amplitude bound	10
1.7.5	6.5 Bootstrap argument	13
1.7.6	6.6 Unconditional regularity via wavevector variance contraction	14
1.8	7. Discussion	18
1.8.1	7.1 Relation to known results	18
1.8.2	7.2 Why the identity is specific to SQG	19
1.8.3	7.3 Computational implications	19
1.8.4	7.4 Comparison with prior conditional regularity results	19
1.9	8. Closed-Form Curvature Correction to the Identity	20
1.9.1	8.1 Single-mode derivation	20
1.9.2	8.2 Generalization to arbitrary front shapes	21
1.9.3	8.3 Four independent verifications	21
1.9.4	8.4 Comparison to 2D Euler	22
1.9.5	8.5 Resolution of the “effective exponent” confusion	22
1.9.6	8.6 Implementation note	23
1.10	9. The Heartbeat Mechanism and Local Energy Budget	23
1.10.1	9.1 Maximum conditions at the gradient maximum	23
1.10.2	9.2 Tangential stretching at gradient growth points	23
1.10.3	9.3 The heartbeat mechanism	24
1.10.4	9.4 Local strain energy budget	24
1.10.5	9.5 The battery argument	26
1.10.6	9.5.1 Local spectral strain energy (the closing measurement)	26

1.10.7	9.5.2 The ion-gate mechanism (dynamical angular damping)	28
1.10.8	9.6 Direct curvature control at the gradient maximum	30
1.10.9	9.7 Proof summary	33
1.10.10	9.8 Reformulation: the thermostat inequality	34
1.10.11	9.8.5 Resolution convergence and the push–pull decomposition	37
1.11	10. Discussion and Open Extensions	38
1.11.1	10.1 Numerical certification of the material max principle	38
1.11.2	10.2 Extension to the generalized SQG family	38
1.11.3	10.3 Implications for ocean reconstruction	39
1.12	References	39
1.13	Appendix: Numerical Methods	39

1 The Shear-Vorticity Identity and Spectral Concentration in SQG Front Dynamics

Bryan Sanchez (*Independent*)

Status (revised 2026-05-29). §2 (the shear–vorticity identity) and §8 (the closed-form curvature correction) are correct, verified, and stand. **The regularity material (§§3–6, 9: selection rule, “material maximum principle”, Theorem 2, “three converging routes”) is WITHDRAWN.** Audit found the Lean “conditional Theorem 2” to be circular (the hypotheses (H-strain)/(H-bdry)/(H- α) are logically vacuous Lean structures and the §9 geometric core is True-stubbed; the only real antecedent is a uniform H^1 bound $\approx$ the conclusion), and the §9 material-maximum-principle argument to be contradicted by this paper’s own numerics ($\text{corr}(f, -) = 0.44$ defeats the free-derivative step $_max \sim G$; the depletion mechanism is provably one order short, needing $_ \sim G^{\{1.2\}}$ blow-up of the $dG/dt = G|nSn|$ bound). The genuine obstruction is the log-criticality $\|S\|_\infty \sim G \log G$. Read §§3–6, 9 as heuristic mechanism, not proof. Full account + salvage scope: NoetherSolve docs/findings/ `sqg_material_max_principle_falsified_2026_05_29.md` and `sqg_salvage_identity_and_curvature_correction_2026_05_29.md`.

1.1 Abstract

I report an algebraic identity for the inviscid Surface Quasi-Geostrophic (SQG) equation: the combination $S_{nt} - \omega/2$ — shear strain minus half the vorticity — vanishes identically for any one-dimensional front configuration. In Fourier space, this combination has the multiplier $|k| \sin^2 \varphi$, where φ is the angle from the front-normal axis. The identity has three consequences for the regularity problem. First, the curvature forcing $F_{\text{ext}} = \partial(S_{nt} - \omega/2)/\partial s$ in the level-set curvature budget vanishes for straight fronts, reducing the budget to pure kinematic straightening. Second, the material wavevector rotation near a compressive front takes the form $d\varphi/dt = 2(nSn)\varphi$ (no constant forcing), driving spectral concentration toward the front-normal axis. Third, the normal strain nSn at the gradient maximum — the quantity controlling potential blowup — is suppressed by the same angular concentration that the identity creates. Pseudospectral simulations at $N = 512$ confirm all three predictions: during front sharpening ($G = 12 \rightarrow 44$), the angular spread narrows as $\Delta\varphi_{50} \sim G^{-1.0}$, the enstrophy concentration ratio increases as $\rho(10^\circ) \sim G^{0.38}$ ($R^2 = 0.98$), and $|nSn(x^*)|$ stays bounded in $[0.17, 0.97]$ across all snapshots — $O(1)$, G -independent — with a weak scaling exponent of -0.15 ± 0.05 inside the clean sharpening interval (both Eulerian and fit-window protocols agree at this precision; see §9.5). I further derive a closed-form expression for the residual $S_{nt} - \omega/2$ on a front with small perturbation $y_f(x) = \eta(x)$: $\text{residual}(x) = \mathcal{F}^{-1}[-k^2 \cdot I(|k|\delta) \cdot \hat{\eta}(k)](x)$, where $I(q) =$

$\int_0^\infty u/[\sinh(\pi u/2)\sqrt{q^2+u^2}] du$ has small- q expansion $I(q) = -(2/\pi) \ln q + 0.5951 + O(q^2)$. This formula is verified to <0.2% amplitude accuracy and >0.9995 correlation across four independent tests (analytic, static arbitrary shapes, dynamical evolution, independent clean-room solver).

Keywords: Surface Quasi-Geostrophic equation, finite-time singularity, spectral concentration, vorticity-strain cancellation, front dynamics

1.2 1. Introduction

The inviscid Surface Quasi-Geostrophic (SQG) equation,

$$\partial_t \theta + \mathbf{u} \cdot \nabla \theta = 0, \quad \mathbf{u} = \nabla^\perp (-\Delta)^{-1/2} \theta, \quad (1)$$

on the torus $\mathbb{T}^2$, shares critical Sobolev scaling with the 3D incompressible Navier-Stokes equations. Whether smooth solutions persist globally or develop finite-time singularities is one of the central open problems in mathematical fluid dynamics [1, 2]. The Beale-Kato-Majda analog states: smooth solutions persist on $[0, T)$ if and only if $\int_0^T \|\nabla \theta\|_\infty dt < \infty$ [3].

Potential singularity formation proceeds through front sharpening: thin regions where $|\nabla \theta|$ concentrates, with the gradient maximum $G(t) = \|\nabla \theta(\cdot, t)\|_\infty$ growing in time. At the gradient maximum $x^*(t)$, the growth rate is

$$\frac{dG}{dt} = G |nSn(x^*)|, \quad (2)$$

where $nSn = \hat{n} \cdot S \cdot \hat{n}$ is the normal strain (compressive, $nSn < 0$, for sharpening fronts). The regularity question reduces to: *how fast can $|nSn|$ grow relative to G ?*

The Calderón-Zygmund (CZ) bound gives $|nSn| \leq CG$ (the strain is bounded by the gradient). Any improvement — $|nSn| = o(G)$ as $G \rightarrow \infty$ — would resolve the problem. The selection rule (parity cancellation of the Riesz kernel at the gradient maximum) gives $|nSn| \leq C_1 G \kappa^2 \delta^2$, where κ is the level-set curvature and $\delta = \|\theta\|_\infty / G$ is the front width [4]. This improves on CZ when $\kappa \delta \ll 1$, but controlling the curvature evolution has remained the obstacle.

In this paper, I identify an algebraic identity that provides the missing mechanism.

Lagrangian framing. The argument developed in §9 operates on a *material* segment of the front — a Lagrangian object — rather than on the Eulerian gradient-maximum location. This is thematically consistent with recent results on the solution-map regularity of generalized SQG: the Lagrangian solution map is real analytic, while the Eulerian map is nowhere locally uniformly continuous in Sobolev topology (arXiv:2603.12944, 2026). The regularity argument in this paper relies on this “well-behaved” side of the dichotomy: the material segment expands under incompressibility, and the curvature maximum principle (Prop 9.9) is a Lagrangian statement on an expanding domain.

Scope of claims (reader’s guide). **Theorem 1** (shear-vorticity identity, §2) is proved unconditionally and machine-verified in Lean 4 + mathlib in the companion repository. **Proposition 6.1** (selection rule, §6.1) is stated with a parity-based argument at formal Taylor-expansion level (not at theorem-grade rigor); the Lean library verifies the universal per-mode CZ precursor (`sqg_selection_rule_bound` in `SqgIdentity/Basic.lean`), while the pointwise $\kappa^2 \delta^2$ refinement at the gradient maximum — which is

what §9.6.3 needs — is the parity-based claim of Proposition 6.1. **Theorem 2** (SQG regularity, §9.6.3) — **WITHDRAWN; see the Status banner** — was stated as conditional on two explicit hypotheses, labeled (H-strain) and (H-bdry), which concern the non-degeneracy of normal strain at the tracked curvature maximum and the boundedness of curvature on the boundary of a material segment, respectively; §9.8 provided a reformulation (H- α) + (B-spec). The 2026-05-29 audit found the conditional chain circular and the §9 material-maximum-principle mechanism contradicted by this paper's own numerics, so §§3–6 and 9 are retained for the record and should be read as heuristic mechanism, not proof. The paper's standing contribution is the identity (§2) and the closed-form curvature correction (§8), together with an explicit accounting of what is rigorous, what is parity-heuristic, and what is withdrawn.

1.3 2. The Shear-Vorticity Identity

1.3.1 2.1 Statement

Theorem 1 (Shear-vorticity identity). *For the SQG velocity field $\mathbf{u} = \nabla^\perp(-\Delta)^{-1/2}\theta$ on $\mathbb{T}^2$, the combination of shear strain and vorticity*

$$\mathcal{F}[S_{nt} - \frac{1}{2}\omega](k) = |k| \sin^2 \varphi_k \cdot \hat{\theta}(k), \quad (3)$$

where φ_k is the angle between k and the unit normal $\hat{n}$, and $S_{nt} = \hat{n} \cdot S \cdot \hat{t}$ with $\hat{t} = \hat{n}^\perp$.

In particular, for any one-dimensional configuration $\theta = f(x \cdot \hat{n})$:

$$S_{nt}(x) - \frac{1}{2}\omega(x) = 0 \quad \text{for all } x \in \mathbb{T}^2. \quad (4)$$

1.3.2 2.2 Proof

The SQG velocity in Fourier space:

$$\hat{u}_1(k) = -ik_2 \hat{\theta}/|k|, \quad \hat{u}_2(k) = ik_1 \hat{\theta}/|k|. \quad (5)$$

In the frame (x_n, x_t) aligned with $\hat{n} = \hat{e}_1$, $\hat{t} = \hat{e}_2$ (with $\hat{t}$ obtained by CCW rotation of $\hat{n}$):

Shear strain:

$$\hat{S}_{nt} = \hat{S}_{12} = \frac{1}{2}(ik_2 \hat{u}_1 + ik_1 \hat{u}_2) = \frac{1}{2} \left(\frac{k_2^2}{|k|} - \frac{k_1^2}{|k|} \right) \hat{\theta} = -\frac{k_1^2 - k_2^2}{2|k|} \hat{\theta}. \quad (6)$$

Vorticity:

$$\hat{\omega} = ik_1 \hat{u}_2 - ik_2 \hat{u}_1 = - \left(\frac{k_1^2}{|k|} + \frac{k_2^2}{|k|} \right) \hat{\theta} = -|k| \hat{\theta}. \quad (7)$$

(Note: $ik_1 \cdot (ik_1 \hat{\theta}/|k|) = -k_1^2 \hat{\theta}/|k|$ and $-ik_2 \cdot (-ik_2 \hat{\theta}/|k|) = -k_2^2 \hat{\theta}/|k|$; both contributions are negative, giving $\hat{\omega} = -|k| \hat{\theta}$, equivalent to $\omega = -(-\Delta)^{1/2} \theta$.)

Difference:

$$\hat{S}_{nt} - \frac{1}{2}\hat{\omega} = -\frac{k_1^2 - k_2^2}{2|k|} \hat{\theta} + \frac{|k|}{2} \hat{\theta} = \frac{|k|^2 - (k_1^2 - k_2^2)}{2|k|} \hat{\theta} = \frac{k_2^2}{|k|} \hat{\theta} = |k| \sin^2 \varphi_k \cdot \hat{\theta}. \quad (8)$$

For a one-dimensional front $\theta = f(x_1)$: $\hat{\theta}(k)$ is supported on $k_2 = 0$, so $\sin^2 \varphi_k = 0$. $\square$

1.3.3 2.3 Physical interpretation

The identity states that along a straight front, the shear strain exactly equals half the vorticity: $S_{nt} = +\omega/2$. Physically, the Riesz velocity field creates a shear pattern whose off-diagonal strain component is locked to the vorticity with the same sign.

The multiplier $|k| \sin^2 \varphi$ vanishes on the k_n axis and grows quadratically with the angular deviation φ . For a front with spectral angular spread $\psi \ll 1$, the residual scales as:

$$|S_{nt} - \frac{1}{2}\omega| \lesssim \psi^2 \cdot \|(-\Delta)^{1/2}\theta\|_\infty \lesssim \psi^2 G. \quad (9)$$

1.4 3. Consequences for the Curvature Budget

1.4.1 3.1 Level-set curvature evolution

For a material curve (level set of θ) in 2D incompressible flow, the signed curvature κ evolves as [5]:

$$\frac{D\kappa}{Dt} = (nSn)\kappa + \frac{\partial}{\partial s} \left(S_{nt} - \frac{1}{2}\omega \right). \quad (10)$$

The first term is the *kinematic straightening*: since $nSn < 0$ at a compressive front, $(nSn)\kappa < 0$ for $\kappa > 0$. This is the 2D incompressibility effect: normal compression creates tangential extension, which stretches and straightens the level set.

The second term $F_{\text{ext}} = \partial(S_{nt} - \omega/2)/\partial s$ is the *curvature forcing* from shear variation along the front.

1.4.2 3.2 Application of the identity

By Theorem 1, $S_{nt} - \omega/2 = 0$ for straight fronts. Therefore $F_{\text{ext}} = 0$ for straight fronts, and the curvature budget reduces to:

$$\frac{D\kappa}{Dt} = (nSn)\kappa \quad (\text{straight front}). \quad (11)$$

This is *pure straightening*: $\kappa \rightarrow 0$ exponentially at rate $|nSn|$.

For a curved front with curvature κ and spectral angular spread $\psi \approx \kappa \delta$, the residual forcing is:

$$|F_{\text{ext}}| \lesssim \frac{\partial}{\partial s} (\psi^2 G) \sim \kappa^2 \delta G = \kappa^2 A, \quad (12)$$

where $A = \|\theta\|_\infty$ and $\delta = A/G$. Crucially, this is *independent of G*.

1.5 4. Wavevector Rotation and Spectral Concentration

1.5.1 4.1 Material gradient dynamics

For a passive scalar advected by an incompressible flow, the material gradient evolves exactly as [6]:

$$\frac{Dk_i}{Dt} = -\frac{\partial u_j}{\partial x_i} k_j. \quad (13)$$

Near a front with compressive strain $nSn < 0$, the angular deviation φ of the local wavevector from the k_n axis satisfies:

$$\frac{d\varphi}{dt} = -(S_{nt} - \frac{1}{2}\omega) + 2(nSn)\varphi + O(\varphi^2). \quad (14)$$

1.5.2 4.2 The identity eliminates the forcing

For a one-dimensional front, the identity (Theorem 1) kills the constant forcing term $S_{nt} - \omega/2 = 0$. The angular dynamics reduce to:

$$\frac{d\varphi}{dt} = 2(nSn)\varphi, \quad (15)$$

which is *pure concentration*: since $nSn < 0$, every wavevector rotates toward the k_n axis at rate $2|nSn|$.

Combined with the gradient growth $dG/dt = G|nSn|$:

$$\psi(t) = \psi_0 \left(\frac{G_0}{G(t)} \right)^2 \quad (\text{pure rotation prediction}). \quad (16)$$

In practice, nonlinear spectral energy creation at non-zero φ (from curvature and front interactions) slows the concentration rate. The measured scaling is $\psi \sim G^{-p}$ with $p \approx 1.0$ (§5), compared to the $p = 2$ prediction from pure rotation.

1.5.3 4.3 Self-reinforcing cycle

The identity creates a self-reinforcing mechanism:

1. **Concentration** ($\psi \rightarrow 0$): wavevector rotation (15) drives the spectral support toward k_n .
2. **Identity residual** ($S_{nt} - \omega/2 \rightarrow 0$): Theorem 1 with $\sin^2 \varphi \rightarrow 0$.
3. **Curvature forcing** ($F_{\text{ext}} \rightarrow 0$): the curvature budget (10) becomes pure straightening.
4. **Straightening** ($\kappa \rightarrow 0$): level-set curvature decays exponentially.
5. **Selection rule** ($|nSn| \rightarrow 0$): the parity-based bound $|nSn| \leq C_1 G \kappa^2 \delta^2 \rightarrow 0$.
6. **Return to 1**: with smaller $|nSn|$, the gradient growth $dG/dt = G|nSn|$ decelerates, but concentration continues.

Each step reinforces the others. The fixed point of this cycle is the one-dimensional front: $\kappa = 0$, $\psi = 0$, $S_{nt} - \omega/2 = 0$, $F_{\text{ext}} = 0$, $nSn = 0$ (from the selection rule with $\kappa = 0$).

1.6 5. Numerical Verification

1.6.1 5.1 Setup

Pseudospectral simulations of inviscid SQG (1) on $\mathbb{T}^2 = [0, 2\pi]^2$ with $N = 256$ and $N = 512$ grid points per dimension. Fourth-order Runge-Kutta time integration, 2/3-dealiasing, $\Delta t = 5 \times 10^{-4}$ ($N = 512$; see Appendix for the general choice). Initial condition: multimode $\theta_0 = \sin x \sin y + 0.5 \cos 3x \sin 2y + 0.3 \sin 2x \cos y$.

1.6.2 5.2 Identity verification

At $N = 256$, $G = 12.5$ (clean sharpening phase, $t = 4.2$), I trace the θ -level set through x^* and sample S_{nt} , $\omega/2$, and their difference along the front.

Quantity	Value at x^*	Range along front
S_{nt}	+0.453	[+0.38, +0.55]
$\omega/2$	+0.447	[+0.38, +0.48]
$S_{nt} - \omega/2$	+0.006	[-0.13, +0.03]
F_{ext}	+2.19	[-4.6, +5.5]

The residual $|S_{nt} - \omega/2| \leq 0.08$ along the entire front, compared to individual terms of magnitude ~ 0.45 . The cancellation exceeds 82%. (Both S_{nt} and $\omega/2$ carry the same sign at a given front point — the identity says they are equal to leading order.)

1.6.3 5.3 Spectral concentration

At $N = 512$ during clean sharpening ($G : 7.4 \rightarrow 44.2$, $t : 3.4 \rightarrow 6.5$):

Quantity	Scaling	R^2
$\Delta \varphi_{50}$ (angular half-width)	$\sim G^{-1.0}$	0.95
$\rho(10^\circ)$ (enstrophy within 10° of k_n)	$\sim G^{+0.38}$	0.98
$\ nSn\ /G$	$\sim G^{-1.17}$	—
$\ nSn\ $	$\sim G^{-0.17}$	—

The normal strain $|nSn|$ *decreases* with G : the gradient maximum decelerates as the front sharpens. This is the spectral concentration mechanism in action.

1.6.4 5.4 Local amplitude of $f = S_{nt} - \omega/2$ near x^*

The critical test: how does the amplitude of f in a neighborhood of x^* scale with G ? I measure the maximum $|f|$ within 2δ of x^* along the front (“local amplitude”) at $N = 512$:

G	local amplitude	local amp / A	max $ f $ / G	Phase
11.6	0.106	0.080	0.016	sharpening
17.4	0.106	0.080	0.013	sharpening
25.7	0.121	0.091	0.011	sharpening

G	local amplitude	local amp / A	$\max f / G$	Phase
31.9	0.144	0.108	0.009	sharpening
37.5	0.106	0.079	0.009	sharpening
42.1	0.275	0.206	0.020	pre-cascade
44.2	3.071	2.303	0.092	cascade

During clean sharpening ($G = 11.6 \rightarrow 37.5$, a $3.2\times$ increase): **the local amplitude stays at** $(0.08 \pm 0.01)A$, independent of G . The ratio $\max |f|/G$ decreases monotonically from 0.016 to 0.009 — the identity becomes *more exact* as the front sharpens.

At the pre-cascade transition ($G \approx 42$), the local amplitude begins to grow. At cascade onset ($G \approx 44$), it exceeds $2A$ — the identity breaks as the front geometry becomes complex.

This establishes the key bound: the local amplitude of $f = S_{nt} - \omega/2$ near x^* is $O(A)$, not $O(G)$. Since f passes through zero at x^* (from the identity) with local oscillation amplitude $\approx 0.1A$ and half-wavelength $\approx \delta$:

$$|F_{\text{ext}}(x^*)| = \left| \frac{\partial f}{\partial s} \right|_{x^*} \approx \frac{0.1A}{\delta} = 0.1G. \quad (5.4.a)$$

This confirms $F_{\text{ext}} \leq CG$ with $C \approx 0.1$ during the sharpening phase.

Physical explanation. The local amplitude is bounded by $A = \|\theta\|_\infty$ because: (i) at x^* , $f = 0$ by the identity (the gradient maximum selects the straightest part of the front); (ii) the oscillation near x^* is driven by the far-field contribution (large-scale θ structure, other fronts), which is bounded by $\|\theta\|_\infty$; (iii) the near-field contribution is parity-suppressed (the Riesz kernel’s odd symmetry cancels the leading terms). The far-field kernel decays as $1/|y|^2$, giving a convergent integral $\lesssim A$.

The far-field contribution to nSn itself (from θ at distance $> L/4$ from x^*) is bounded: $|nSn_{\text{far}}| \leq 0.18$ across all snapshots, independent of G .

1.6.5 5.5 Two dynamical regimes

Regime	G range	$\ nSn\ /G$	κ	$\rho(10^\circ)$	Mechanism
Clean sharpening	$12 \rightarrow 44$	$0.07 \rightarrow 0.005$	$0.03\text{--}1.1$	$0.09 \rightarrow 0.53$	Identity + concentration
Turbulent cascade	$48 \rightarrow 96$	≤ 0.10	$15\text{--}55$	$0.08\text{--}0.16$	KH instability disrupts

In the sharpening regime, the self-reinforcing cycle (§4.3) operates. In the turbulent regime, multiple interacting fronts undergo Kelvin-Helmholtz cascade fragmentation, keeping $|nSn|/G$ bounded.

1.7 6. Implications for SQG Regularity

1.7.1 6.1 The selection rule

Proposition 6.1 (Selection rule, parity-based refinement). *Let $\theta \in C^\infty(\mathbb{T}^2)$ and let x^* be a gradient maximum of θ , with level-set curvature κ at x^* and front width $\delta = \|\theta\|_\infty / \|\nabla \theta\|_\infty$. Assume the gradient-maximum conditions $\theta_{nn}(x^*) = \theta_{nt}(x^*) = 0$ of Proposition 9.1. Then the near-field contribution to the normal strain satisfies*

$$|nSn_{\text{near}}(x^*)| \leq C \kappa^2 \delta^2 \|\nabla \theta\|_\infty = C \kappa^2 \|\theta\|_\infty^2 / \|\nabla \theta\|_\infty, \quad (\text{S})$$

where C is a constant depending only on the CZ kernel, $\|\Theta'''\|_\infty$ of the 1-D profile, and the cutoff radius R ; and nSn_{near} denotes the contribution to $nSn(x^*)$ from the near-field shell $|y - x^*| < R\delta$ of the Riesz velocity's kernel (Lemma 6.2).

Argument. Near x^* the front is one-dimensional up to curvature $O(\kappa)$. The Riesz velocity's CZ kernel K_f is even-even in (n, t) ; three of the four parity sectors of the non-1D perturbation $\Delta\theta = \theta - \theta_{1D}$ integrate to zero exactly. The surviving (*ee*) sector's first-order-in-curvature contribution has angular average zero on each radial shell (by the vanishing angular mean of the CZ kernel's (*ee*) part), so it does not contribute at that order. The leading surviving contribution is at second order in curvature, using $\theta_{nn}(x^*) = \theta_{nt}(x^*) = 0$ (Prop 9.1) to remove the n^1 coefficient of $\Theta'(n)$, giving the $\kappa^2 \delta^2$ bound. Full derivation with constants: Lemmas 6.2, 6.3. Cf. Córdoba 1998 [4] for the unconditional $O(\kappa A)$ form.

Remark (rigor level and relation to Lean). Proposition 6.1 is stated as an *unconditional bound* on a specific near-field functional (not a regularity result). Its proof (Lemma 6.2 below) operates at the level of formal Taylor expansion plus parity-sector cancellation, with constants tracked dimensionally rather than sharply; it is tighter than the rigorous Córdoba [4] $O(\kappa A)$ bound and is the ingredient that makes bounded curvature suffice in §9.6.3. **Proposition 6.1 itself is not machine-verified:** the companion Lean library verifies the *universal per-mode CZ bound* $\|S_{nt} - \omega/2(k)\| \leq |k| \|\hat{\theta}(k)\|$ (see `sqg_selection_rule_bound` in `SqgIdentity/Basic.lean`), which Proposition 6.1 refines at the gradient maximum via the parity argument. The `MaterialMaxPrinciple` chain in `SqgIdentity/RieszTorus.lean` §14 takes the *integrated* $\dot{H}^1$ form of the per-mode bound as one of its spectral inputs, not the pointwise $\kappa^2 \delta^2$ refinement.

1.7.2 6.2 Self-consistent bound

The curvature budget (10) with $F_{\text{ext}} \leq C_2 G$ (verified in §5.4 during sharpening) gives the equilibrium curvature $\kappa_{\text{eq}} = F_{\text{ext}}/|nSn|$. Combined with the selection rule (S) of Proposition 6.1:

$$|nSn|^3 \leq \frac{C_1 C_2^2 A^2}{4} G, \quad \text{i.e.,} \quad |nSn| \leq C G^{1/3}. \quad (17)$$

1.7.3 6.3 Cascade argument

The Kelvin-Helmholtz instability of a front with width $\delta = A/G$ and velocity jump $\sim A$ has growth rate $\gamma_{\text{KH}} \sim G$. The self-consistent sharpening rate is $|nSn| \leq C G^{1/3}$. The ratio:

$$\frac{\gamma_{\text{KH}}}{|nSn|} \sim G^{2/3} \rightarrow \infty \quad \text{as } G \rightarrow \infty. \quad (18)$$

The KH instability accumulates $G^{2/3}$ e-foldings during each sharpening episode — infinitely more than needed to disrupt the front. After disruption, the gradient maximum does not increase (the front fragments into smaller-amplitude structures). The BKM integral $\int G dt$ remains finite on any bounded time interval.

1.7.4 6.4 The local amplitude bound

The self-consistent argument (§6.2) and the cascade (§6.3) together prove regularity, *conditional on* $F_{\text{ext}}(x^*) \leq CG$. Section 5.4 establishes this empirically: the local amplitude of $f = S_{nt} - \omega/2$ near x^* satisfies

$$\max_{|s| < 2\delta} |f(x^* + s\hat{t})| \leq (0.1 \pm 0.02) A \quad (19)$$

across the entire clean sharpening phase at $N = 512$ (G from 11.6 to 37.5). Since $f(x^*) \approx 0$ and the half-wavelength of the oscillation is $\sim \delta$, the derivative bound (13) follows: $F_{\text{ext}}(x^*) \leq 0.1 G$.

The bound (19) has a clear structural explanation:

1. **Near-field parity.** The Riesz kernel's odd symmetry, combined with the gradient maximum conditions $\theta_{nn} = \theta_{nt} = 0$ and the identity $f = 0$ for 1D fronts, suppresses the near-field contribution to f near x^* by two orders (quadratic in the angular deviation $\kappa\delta$).
2. **Far-field boundedness.** The far-field contribution to f at any point on the front is bounded by CA (from $\|\theta\|_\infty$ conservation and the $1/|y|^2$ kernel decay of the second-order operator). This is G -independent.
3. **Gradient-max selection.** The gradient maximum x^* sits at the straightest part of the front, where the near-field contribution is minimized and the far-field dominates. The local amplitude is set by the far field: $O(A)$.

The analytical proof of (19) rests on three lemmas, proved below.

Lemma 6.1 (Far-field bound). *Let θ be a smooth SQG solution on $\mathbb{T}^2$ with $A := \|\theta\|_\infty$, and let $f = S_{nt} - \omega/2 = \partial_t^2(-\Delta)^{-1/2}\theta$. Fix a point $x_0 \in \mathbb{T}^2$ and a length scale $R_0 \in (0, L/4]$ independent of $G = \|\nabla\theta\|_\infty$. Let $\chi \in C_c^\infty(\mathbb{T}^2)$ be a cutoff with $\chi \equiv 1$ on $B(x_0, R_0/2)$, $\text{supp}(\chi) \subset B(x_0, R_0)$, and $\|\chi\|_{C^m} \leq C_m R_0^{-m}$. Write $\theta = \theta_{\text{near}} + \theta_{\text{far}}$ with $\theta_{\text{near}} := \chi\theta$ and $\theta_{\text{far}} := (1 - \chi)\theta$. Then for every integer $k \geq 0$,*

$$\|\nabla^k [(-\Delta)^{-1/2}\theta_{\text{far}}]\|_{L^\infty(B(x_0, R_0/4))} \leq C_{k, R_0} A, \quad (6.1.a)$$

with C_{k, R_0} depending only on k , R_0 , and the torus period. In particular,

$$|f_{\text{far}}(x_0)| \leq C_{R_0} A, \quad |nS_{n\text{far}}(x_0)| \leq C_{R_0} A, \quad (6.1.b)$$

uniformly in G .

Proof. On $\mathbb{T}^2 = \mathbb{R}^2/(2\pi\mathbb{Z})^2$, the Riesz potential of order 1 has periodic kernel $G_{1/2}(z) = c_0|z|^{-1} + r_{\text{per}}(z)$, where $c_0 = 1/(2\pi)$ and $r_{\text{per}} \in C^\infty(\mathbb{T}^2)$ is the lattice correction (real-analytic off the origin, with exponentially decaying Fourier coefficients). For $x \in B(x_0, R_0/4)$ and $y \in \text{supp}(\theta_{\text{far}}) \subset \{|y - x_0| \geq R_0/2\}$, we have $|x - y| \geq R_0/4$, so $G_{1/2}(x - y)$ is C^∞ jointly in (x, y) with derivatives bounded by $C_k R_0^{-k-1}$. Differentiation under the integral gives

$$|\nabla_x^k [(-\Delta)^{-1/2} \theta_{\text{far}}](x)| = \left| \int \nabla_x^k G_{1/2}(x-y) \theta_{\text{far}}(y) dy \right| \leq \|\theta\|_\infty \cdot \|\nabla^k G_{1/2}\|_{L^1(|z| \geq R_0/4)} \leq C_{k,R_0} A,$$

proving (6.1.a). Since $f = \partial_t^2(-\Delta)^{-1/2} \theta$ and the components of S are each of the form $\partial_i \partial_j (-\Delta)^{-1/2} \theta$ composed with a bounded multiplier, applying (6.1.a) with $k = 2$ yields (6.1.b). $\square$

Remark 6.1.1. Unlike the Calderón-Zygmund bound $|nSn| \leq CG$, the far-field bound (6.1.b) is *independent of G* : the smoothing action of $(-\Delta)^{-1/2}$ on sources separated from the query point converts the source L^∞ norm A into a pointwise bound on arbitrary-order derivatives of $\psi = (-\Delta)^{-1/2} \theta$ without invoking the gradient norm. The fixed scale R_0 (as opposed to a G -dependent scale like $R\delta$) is what makes the argument elementary.

Lemma 6.2 (Near-field parity suppression). *At a gradient maximum x^* , where the gradient-maximum conditions $\theta_{nn}(x^*) = \theta_{nt}(x^*) = 0$ of Proposition 9.1 hold, the near-field contribution*

$$f_{\text{near}}(x^*) = \int_{|y-x^*| < R\delta} K_f(x^* - y) \theta(y) dy$$

satisfies $|f_{\text{near}}(x^)| \leq C\kappa^2 \delta^2 G = C\kappa^2 A^2/G$, where κ is the level-set curvature at x^* and K_f is the CZ kernel for $f = \partial_t^2(-\Delta)^{-1/2} \theta$.*

Proof. Work in local coordinates (n, t) centered at x^* . The CZ kernel $K_f(n, t) = c_f(t^2 - n^2)/(n^2 + t^2)^{5/2} +$ lattice correction is even in n , even in t (it is the real part of ∂_t^2 applied to $1/|z|$, which has even-even symmetry in the two axes). The windowed field $W_\theta(n, t) := \chi_{B(R\delta)} \theta(x^* + n\hat{n} + t\hat{t})$ expands in parity sectors as

$$W_\theta = W^{(ee)} + W^{(eo)} + W^{(oe)} + W^{(oo)},$$

where $W^{(ij)}$ is even/odd in n and even/odd in t respectively. Because K_f is even-even in (n, t) , only the $W^{(ee)}$ sector contributes to $f_{\text{near}}(x^*) = (K_f * W_\theta)(0)$; the other three sectors integrate to zero exactly.

Write the 1-dimensional profile $\theta_{1D}(n) = \Theta(n)$ with $\Theta'(0) = G$, $\Theta''(0) = \theta_{nn}(x^*) = 0$. By Theorem 1, the 1-dimensional sector of W_θ (which is $\theta_{1D}(n)$, independent of t , and is even in t trivially, and has parity in n determined by Θ) contributes zero to $f_{\text{near}}(x^*)$: indeed, for a straight front the full integrand convolves to zero by the identity. So in $W^{(ee)}$ we can subtract θ_{1D} without changing f_{near} .

Taylor-expand the curved front. With signed distance $d(n, t) = n - \frac{1}{2}\kappa t^2 + O(\kappa^2 t^3 n + \kappa t^4)$,

$$\theta(n, t) = \Theta(d(n, t)) = \Theta(n) - \frac{1}{2}\kappa t^2 \Theta'(n) + \frac{1}{8}\kappa^2 t^4 \Theta''(n) + O(\kappa^2 t^3 n \Theta') + O(\kappa^3).$$

Let $\Delta\theta := \theta - \Theta$ be the curvature-induced deviation. We extract its (ee) component $\Delta\theta^{(ee)}$:

- The term $-\frac{1}{2}\kappa t^2 \Theta'(n)$ is *even in t* , but the factor $\Theta'(n)$ has parity structure $\Theta'(0) + \Theta''(0)n + \frac{1}{2}\Theta'''(0)n^2 + \dots = G + \frac{1}{2}\Theta'''(0)n^2 + \dots$ at x^* , using $\Theta''(0) = \theta_{nn}(x^*) = 0$. So $\Theta'(n)$ is *even in n* to leading order (since the n^1 coefficient vanishes by the gradient-maximum condition). Hence the $\frac{1}{2}\kappa t^2 \Theta'(n)$ piece is (ee) at $O(\kappa)$.
- However, this $O(\kappa)$ (ee) piece is precisely the curvature correction to the 1D Riesz potential Ψ_0 of Lemma 6.3 (the Ψ_1 term), whose tangential second derivative at $d = 0$ gives $-\kappa \Psi_0'(0)$. This leading piece *does* appear in $f(x^*)$ — it is the $-c\kappa A$ term of Lemma 6.3, not the near-field f_{near} under the CZ cutoff $|y - x^*| < R\delta$. With the cutoff, it is a *macroscale* (far-field) contribution from the curved profile at distances $|t| \sim 1/\sqrt{\kappa} \gg \delta$; at distances $|t| < R\delta$ with $\kappa\delta \ll 1$, the factor $\kappa t^2 < \kappa R^2 \delta^2$ is $O(\kappa^2 \delta^2)$ after integrating against the CZ kernel, as we show next.

- The $O(\kappa^2)$ remainders in the Taylor expansion are (ee) as well, and contribute at higher order.

For the near-field contribution at $|z| < R\delta$: the (ee) part of $\Delta\theta$ satisfies $|\Delta\theta^{(ee)}(n, t)| \leq C\kappa t^2 (G + C' n^2)$ on the disk, and hence the cutoff $|(n, t)| < R\delta$ gives $|\Delta\theta^{(ee)}| \leq C\kappa R^2 \delta^2 G$ pointwise (using $\Theta''(0) = 0$ so the n^2 term is subleading). Convolution with the CZ kernel, noting that $K_f \sim 1/r^3$ and the (ee) sector survives:

$$|f_{\text{near}}(x^*)| \leq C\kappa R^2 \delta^2 G \int_0^{R\delta} \frac{1}{r^3} r dr \cdot (\text{angular factor from } (ee) \text{ sector}).$$

The radial integral $\int_0^{R\delta} r^{-2} dr$ is logarithmically divergent at zero; however, the (ee) angular average of K_f over the full unit circle vanishes by parity (the CZ kernel has zero average over angular shells of a fixed radius), so the divergent near-origin contribution vanishes. The leading non-vanishing contribution then comes from the *quadratic* Taylor correction $-\frac{1}{2}\kappa t^2 \cdot \frac{1}{2}\Theta'''(0)n^2$ (recall the n^1 term in Θ' vanishes at x^*), giving the second (ee) cancellation structure $(n^2 t^2)/r^5$. After angular averaging, the surviving piece is $O(\kappa^2 G/\delta) \cdot (R\delta)^2 = O(\kappa^2 R^2 \delta G) = O(\kappa^2 A R^2)$, which with $\delta = A/G$ is $O(\kappa^2 \delta^2 G)$.

This establishes

$$|f_{\text{near}}(x^*)| \leq C\kappa^2 \delta^2 G = C\kappa^2 A^2/G. \quad (6.2.a)$$

The constant C depends only on R , the 1-D profile Θ through $\|\Theta'''\|_\infty$, and the CZ kernel's angular harmonics; crucially it is independent of G .

The same parity argument applies verbatim to the normal-strain multiplier $m_{nn}(k) = -(|k|/2) \sin(2\varphi_k)$, which is (oo) in (k_n, k_t) , hence dual to (oo) in (n, t) on the physical side; the surviving sector is again the $O(\kappa^2)$ one, giving $|n S n_{\text{near}}(x^*)| \leq C\kappa^2 A^2/G$ as stated in Proposition 6.1. $\square$

Remark 6.2.1 (where the gradient-maximum conditions are essential). The parity argument above *requires* $\theta_{nn}(x^*) = \theta_{nt}(x^*) = 0$ at x^* (Proposition 9.1): without them, the 1-D profile Θ would contain an n^1 term in $\Theta'(n)$, and the (ee) sector would receive an $O(\kappa)$ (rather than $O(\kappa^2)$) contribution. At a generic non-maximum point, only the weaker $O(\kappa A)$ bound of the original Córdoba-type argument holds. This is why Proposition 6.1 is stated at the gradient maximum x^* specifically, and why the bootstrap of §6.5 tracks the interior maximum of curvature on a material segment rather than a generic front point.

Lemma 6.3 (Curvature tracking, with controlled remainder). *Let θ be a smooth SQG solution on a short time interval around a front segment with curvature $\kappa(s)$ parametrized by arclength s , and suppose the front width $\delta = A/G$ is small compared to the torus period. There exist a profile-dependent constant $c = c(\theta_0) \in \mathbb{R}$ and a remainder function $r(s)$ such that*

$$f(s) = S_{nt}(s) - \frac{1}{2}\omega(s) = -c \kappa(s) A + r(s), \quad r(s) = A \kappa^2(s) g(s), \quad (6.3.a)$$

where $g \in C^1$ satisfies the uniform bound

$$\|g\|_{L^\infty} + \|g'\|_{L^\infty} \leq C(\theta_0), \quad (6.3.b)$$

with $C(\theta_0)$ independent of G .

Proof. Write $\psi = (-\Delta)^{-1/2}\theta$. For a curved front with profile $\theta(x) = \Theta(d(x))$, $\Theta'(0) = G$, and signed distance d , the asymptotic expansion of the Riesz layer potential in powers of curvature gives

$$\psi(x) = \Psi_0(d) + \kappa(s) \Psi_1(d) + \kappa^2(s) \Psi_2(d, s) + O(\kappa^3), \quad (6.3.c)$$

where $\Psi_0 = (-\Delta)^{-1/2}\Theta$ is the one-dimensional Riesz potential, Ψ_1 is the universal first curvature correction, and $\Psi_2(d, s)$ carries the arclength dependence at second order. Taking the tangential second derivative on the front ($d = 0$), using $\partial_t d = 0$ and $\partial_t^2 d = -\kappa$:

$$f(s) = \partial_t^2 \psi(s) = -\kappa(s) \Psi'_0(0) + \kappa^2(s) \Psi_2^{(0,0)}(0, s) + O(\kappa^3), \quad (6.3.d)$$

where $\Psi_2^{(0,0)}(0, s)$ denotes the value of Ψ_2 at $d = 0$ as a function of s . Set $c := \Psi'_0(0)/A$, so the leading term is $-c\kappa A$. For the remainder, write the full stream function via the cutoff decomposition of Lemma 6.1: $\psi = \psi_{\text{near}} + \psi_{\text{far}}$ with cutoff scale R_0 chosen as a small constant times the torus period.

- **Near-field contribution to r .** The near-field part ψ_{near} inherits the explicit asymptotic form (6.3.c), and $\Psi_2(0, s)$ is a profile-dependent constant plus slowly-varying geometric corrections whose s -derivative is controlled by the front-scale geometry uniformly in G .
- **Far-field contribution to r .** By Lemma 6.1 with $k = 2, 3$: $\partial_t^2 \psi_{\text{far}}$ and $\partial_s \partial_t^2 \psi_{\text{far}}$ are bounded pointwise on $B(x_0, R_0/4)$ by $C(\theta_0, R_0)A$, independently of G . This contribution is absorbed into $A \kappa^2(s) g(s)$ with $g = O(A^{-1} \kappa^{-2}) \times (CA) = O(1)$ on the range of κ of interest (the interior of the material segment, where κ is separated from 0); at curvature zeros it is absorbed into the $O(\kappa^3)$ term of (6.3.d), which is of the same form with additional κ suppression.

Combining, $r(s) = A\kappa^2(s)g(s)$ with $g, g' \in L^\infty$ uniformly in G , giving (6.3.b). The constant c is $O(1)$ and scales with A , not G , since $(-\Delta)^{-1/2}$ reduces one derivative. The numerical value $c \approx 0.14$ observed in §5.4 is one particular realization. $\square$

Remark 6.3.1 (why the derivative bound is needed). The bound (6.3.b) on g' is not an aesthetic refinement: it is exactly what the curvature-maximum argument of §9 (Prop 9.9) requires. Differentiating (6.3.a) gives $F_{\text{ext}}(s) = -cA\kappa'(s) + 2A\kappa\kappa'g + A\kappa^2g'$; at an interior curvature maximum κ' vanishes but g' need not, and the surviving $A\kappa^2g'$ term must be estimated separately. Without (6.3.b), the chain would invoke a hidden assumption.

Combining Lemmas 6.1–6.3: the local amplitude of f near x^* is at most CA (from the far field) plus $C\kappa A$ (from the near-field curvature). Since κ is bounded by the bootstrap (Proposition below), the total is $O(A)$, establishing (19).

1.7.5 6.5 Bootstrap argument

The empirical relationship $f \approx -c\kappa A$ (§5.4) enables a self-closing bootstrap.

Observation (curvature tracking). Along the front through x^* , the identity residual tracks the local level-set curvature:

$$f(s) = S_{nt}(s) - \frac{1}{2}\omega(s) \approx -c\kappa(s)A, \quad c \approx 0.14 \quad (20)$$

with correlation $\text{corr}(f, -\kappa) = 0.44$ and quantitative agreement $|f| \leq 0.14\kappa A$ throughout the front. This relationship follows from the multiplier structure: $f = \partial_t^2 [(-\Delta)^{-1/2}\theta]$, and the second tangential derivative of a curved front of amplitude A is proportional to κA .

Proposition (Gronwall closure). Assume (20) holds and $nSn < 0$ along the entire front through x^* (the “clean sharpening” condition, verified in all sharpening-phase snapshots). Then $\kappa_{\text{max}}(t)$ — the maximum curvature along the front — is non-increasing, and $|nSn(x^*)|$ is bounded.

Proof. At the point s_{max} where κ achieves its maximum along the front:

- (a) $\partial \kappa / \partial s = 0$ (necessary condition for a maximum).
- (b) $F_{\text{ext}}(s_{\text{max}}) = \partial f / \partial s \approx -cA \partial \kappa / \partial s = 0$ (from (20)).
- (c) The curvature budget (10) at s_{max} :

$$\frac{D\kappa_{\text{max}}}{Dt} \leq (nSn) \kappa_{\text{max}} + F_{\text{ext}}(s_{\text{max}}) = (nSn) \kappa_{\text{max}} \leq 0 \quad (21)$$

since $nSn < 0$ by the sharpening condition.

Therefore $\kappa_{\text{max}}(t) \leq \kappa_{\text{max}}(0)$ for all t in the sharpening phase. The consequences cascade:

$$|f| \leq c \kappa_{\text{max}} A \leq c \kappa_{\text{max}}(0) A = O(A), \quad (22)$$

$$|F_{\text{ext}}(x^*)| \approx \frac{|f|_{\text{local}}}{\delta} \leq \frac{c \kappa_{\text{max}}(0) A}{\delta} = c \kappa_{\text{max}}(0) G, \quad (23)$$

$$|nSn(x^*)| \leq C_1 \kappa_{\text{max}}^2 \delta^2 G \leq C_1 \kappa_{\text{max}}(0)^2 A^2 / G \rightarrow 0 \quad \text{as } G \rightarrow \infty. \quad (24)$$

From (24): $dG/dt = G |nSn| \leq C_1 \kappa_{\text{max}}(0)^2 A^2$, so $G(t) \leq G(0) + C_1 \kappa_{\text{max}}(0)^2 A^2 t$ — *at most linear growth*. The BKM integral $\int_0^T G dt \leq G(0)T + \frac{1}{2} C_1 \kappa_{\text{max}}(0)^2 A^2 T^2 < \infty$ for any finite T . $\square$

Remark (self-reinforcement). The bound improves with time: as κ_{max} decreases (from (21)), the selection rule bound (24) tightens, which reduces $|nSn|$, which accelerates the straightening. The fixed point is the one-dimensional front: $\kappa = 0, f = 0, nSn = 0$. The sharpening phase converges toward this fixed point.

Remark (conditional nature of the Gronwall closure). The Gronwall closure requires $nSn < 0$ on the *entire* front through x^* — the “clean sharpening” condition. Numerical tests (§5) show this condition *never* holds in practice: extensional regions ($nSn > 0$) always exist somewhere on the front, even when $nSn(x^*) < 0$ at the gradient maximum. I therefore require an unconditional argument. The following proposition achieves this by exploiting the wavevector rotation structure directly, bypassing the whole-front condition.

1.7.6 6.6 Unconditional regularity via wavevector variance contraction

The key observation is that the wavevector rotation equation (14)–(15) provides *automatic* spectral concentration at any point of gradient growth, without requiring control of nSn away from x^* .

Lemma 6.4 (Exact wavevector rotation). *For a wavevector $\mathbf{k}$ advected by the material gradient equation $Dk_i/Dt = -(\partial u_j / \partial x_i) k_j$ in a 2D incompressible flow, the angle φ of $\mathbf{k}$ from the front-normal direction $\hat{n}$ evolves as*

$$\frac{d\varphi}{dt} = (nSn) \sin(2\varphi) - S_{nt} \cos(2\varphi) + \frac{\omega}{2}. \quad (25)$$

Proof. In the $(\hat{n}, \hat{t})$ frame, let $\hat{k} = (\cos \varphi, \sin \varphi)$. The velocity gradient decomposes as $\nabla \mathbf{u} = S + \Omega$ where $S = \begin{pmatrix} a & b \\ b & -a \end{pmatrix}$ with $a = nSn, b = S_{nt}$, and $\Omega = \begin{pmatrix} 0 & -\omega/2 \\ \omega/2 & 0 \end{pmatrix}$. The material gradient equation $d\hat{k}/dt = -(\nabla \mathbf{u})^T \hat{k}$ (projected to the angular component) gives:

$$\frac{d\varphi}{dt} = \hat{k}_\perp \cdot [-(\nabla \mathbf{u})^T \hat{k}] = a \sin(2\varphi) - b \cos(2\varphi) + \frac{\omega}{2},$$

which is (25). $\square$

Verification of (25): At $\varphi = 0$ (pure front-normal wavevector, corresponding to a 1D front): $d\varphi/dt = -S_{nt} + \omega/2 = -(S_{nt} - \omega/2) = 0$ by Theorem 1. Consistent: a 1D front stays 1D. For small φ : $\sin(2\varphi) \approx 2\varphi$, $\cos(2\varphi) \approx 1 - 2\varphi^2$, recovering the linearized form (14): $d\varphi/dt \approx -(S_{nt} - \omega/2) + 2(nSn)\varphi$.

Lemma 6.5 (Variance contraction). *At any point x where $nSn(x) < 0$, consider the ensemble of wavevectors $\{\mathbf{k}_\alpha\}$ carried to x along material trajectories, with angles $\{\varphi_\alpha\}$ from $\hat{n}$. The variance $V(t) = \langle (\varphi_\alpha - \bar{\varphi})^2 \rangle$ (where $\bar{\varphi} = \langle \varphi_\alpha \rangle$ is the weighted mean angle) satisfies:*

$$\frac{dV}{dt} = 4(nSn) V. \quad (26)$$

In particular, since $nSn < 0$: $V(t)$ decreases monotonically. Along a material trajectory where the gradient grows from G_0 to $G(t)$:

$$V(t) = V_0 \left(\frac{G_0}{G(t)} \right)^4. \quad (27)$$

Proof. Equation (25) has the form $d\varphi_\alpha/dt = F(t) + 2(nSn)\varphi_\alpha + R(\varphi_\alpha)$, where $F(t) = -S_{nt} + \omega/2$ is independent of α (it depends on the velocity field at x , not on the individual wavevector), and $R(\varphi) = (nSn)[\sin(2\varphi) - 2\varphi] - S_{nt}[\cos(2\varphi) - 1]$ collects the higher-order terms. To leading order in the angular spread $\psi = \sqrt{V}$ (valid when $\psi \ll 1$):

The difference between any two wavevectors evolves as:

$$\frac{d(\varphi_\alpha - \varphi_\beta)}{dt} = 2(nSn)(\varphi_\alpha - \varphi_\beta) + [R(\varphi_\alpha) - R(\varphi_\beta)]. \quad (28)$$

The forcing $F(t)$ cancels exactly — it is the same for all wavevectors at the same point. The remainder $R(\varphi_\alpha) - R(\varphi_\beta)$ is $O(|\varphi_\alpha - \varphi_\beta| \cdot \max(|\varphi_\alpha|, |\varphi_\beta|))$ since $R' = O(\varphi)$.

For the leading-order term: $d(\Delta\varphi)/dt = 2(nSn)\Delta\varphi$, giving $\Delta\varphi(t) = \Delta\varphi(0) \exp(2 \int_0^t nSn d\tau)$. The variance $V = \langle (\Delta\varphi)^2 \rangle$ satisfies:

$$V(t) = V(0) \exp\left(4 \int_0^t nSn d\tau\right). \quad (29)$$

Since $dG/dt = |nSn|G = -nSn \cdot G$ (equation (2), with $nSn < 0$), I have $\int_0^t nSn d\tau = -\int_0^t (dG/G) = -\ln(G/G_0)$. Substituting:

$$V(t) = V_0 e^{-4 \ln(G/G_0)} = V_0 \left(\frac{G_0}{G} \right)^4. \quad (30)$$

The higher-order correction from R : for concentrated spectra with $|\varphi_\alpha| \leq \psi_{\max}$, the nonlinear remainder satisfies $|R(\varphi_\alpha) - R(\varphi_\beta)| \leq C\psi_{\max}|\varphi_\alpha - \varphi_\beta|$ (since $|R'(\varphi)| = |2(nSn)[\cos(2\varphi) - 1] + 2S_{nt}\sin(2\varphi)| \leq C(|nSn| + |S_{nt}|)\varphi \leq CG\varphi$). This contributes a correction of relative order $\psi_{\max}$ to the variance evolution:

$$\frac{dV}{dt} = [4(nSn) + O(G\psi_{\max})] V.$$

For the variance contraction to dominate, I need $4|nSn| \gg CG\psi_{\max}$. Since $|nSn| \leq CG\psi$ (from the multiplier structure, §4), this requires $4CG\psi \gg CG\psi$, which holds with margin. More precisely, the nonlinear term makes the contraction *faster* (since $|\sin(2\varphi)| < 2|\varphi|$ for $\varphi \neq 0$, the exact equation contracts the angular spread MORE than the linearized one). Therefore (30) is an *upper bound* on the variance. $\square$

Remark (factor of 4 in the variance). The exponent 4 in (27) arises from two compounding factors of 2: (a) the material gradient equation $d\hat{k}/dt = -(\nabla \mathbf{u})^T \hat{k}$ projected to the angular component gives the strain contribution to $d\varphi/dt$ as $a \sin(2\varphi)$, where the $\sin(2\varphi)$ (rather than $\sin \varphi$) is the tensor-algebraic signature of a symmetric traceless 2×2 strain in the angular representation; (b) squaring $d\varphi/dt$ to get dV/dt contributes another factor of 2 through the identity $d(\varphi^2)/dt = 2\varphi d\varphi/dt$. The product $2 \times 2 = 4$ is the exponent in $V(t) = V_0(G_0/G)^4$. The proportionality $d\varphi/dt \supset 2a\varphi = 2(nSn)\varphi$ at small φ linearizes the exact ODE (25); the factor of 2 there is the $\sin(2\varphi) \rightarrow 2\varphi$ Taylor expansion. This is the reason spectral concentration outpaces the forcing that would disrupt it: the forcing from the identity residual creates angular deviation at rate proportional to $\psi^2 G$ (from the $|k| \cos^2 \varphi$ multiplier), while the damping removes it at rate $4|nSn|\psi^2$. For $\psi > 0$ in the concentration regime, the damping dominates the source.

Lemma 6.6 (Mean angle bound). *The mean wavevector angle $\bar{\varphi}(t)$ satisfies*

$$|\bar{\varphi}(t)| \leq \frac{|S_{nt} - \omega/2|}{2|nSn|} + O(\psi^2), \quad (31)$$

where $\psi = \sqrt{V}$ is the RMS angular spread.

Proof. The mean angle evolves as $d\bar{\varphi}/dt = -(S_{nt} - \omega/2) + 2(nSn)\bar{\varphi} + O(G\psi^2)$. Setting $d\bar{\varphi}/dt = 0$ at the (attracting) equilibrium gives $\bar{\varphi}_{\text{eq}} = (S_{nt} - \omega/2)/(2|nSn|)$. By the identity residual bound (Lemma 6.2): $|S_{nt} - \omega/2| \leq C\psi^2 G$ (from the $|k| \cos^2 \varphi$ multiplier structure, since only modes with $|\varphi| \gtrsim \psi$ contribute to the residual). With $|nSn| \geq c\psi G$ (from the multiplier $|k| \sin(2\varphi) \sim 2\psi|k|$ for typical modes): $|\bar{\varphi}_{\text{eq}}| \leq C\psi^2 G / (2c\psi G) = C'\psi$. That is, the mean shift is bounded by a constant times the RMS spread. $\square$

Proposition 6.7 (Preliminary conditional regularity — superseded by Theorem 2 of §9.6.3). *For smooth initial data $\theta_0 \in C^\infty(\mathbb{T}^2)$, assume the near-field spectral concentration bound*

$$|nSn_{\text{near}}(x(t))| \leq C_{\text{near}} \psi(t) |\nabla \theta|(x(t)) \quad (\text{PC})$$

holds along every material trajectory of gradient growth, with ψ the RMS angular spread of Lemma 6.5. Then the inviscid SQG equation (1) preserves C^∞ regularity for all time.

Remark 6.7.1 (status of the hypothesis (PC)). The hypothesis (PC) — that material wavevector concentration at $x(t)$ implies a pointwise Fourier-spectral bound on θ_{near} — conflates two distinct objects (a material-frame angular variance and a Fourier-norm concentration) and is not derivable in the form stated. Later sections narrow the gap: §9.6.3 replaces (PC) by the cleaner hypothesis pair (H-strain) + (H-bdry), and §9.8 further consolidates them into the single scalar thermostat inequality (H- α). Proposition 6.7 is kept here to motivate the material-frame variance-contraction argument that §9.6 makes rigorous; its proof below illustrates the Dini-derivative + BKM route that the refined Theorem 2 instantiates.

Proof. Define $G(t) = \|\nabla\theta(\cdot, t)\|_\infty$ and suppose for contradiction that $G(t) \rightarrow \infty$ as $t \rightarrow T^*$ (the BKM criterion). Consider any material point $x(t)$ at which the gradient grows: $d|\nabla\theta|/dt > 0$. At such a point, $nSn < 0$ (from the material gradient equation $d|\nabla\theta|/dt = -nSn \cdot |\nabla\theta|$).

Step 1 (Variance contraction). By Lemma 6.5, the spectral angular variance at $x(t)$ satisfies $V(t) \leq V_0(G_0/G(t))^4$, where V_0 and G_0 refer to the initial state of the material parcel. For smooth initial data, $V_0 \leq C(\theta_0)$ (bounded by initial data). The RMS angular spread:

$$\psi(t) = \sqrt{V(t)} \leq C(\theta_0) \left(\frac{G_0}{G(t)} \right)^2. \quad (32)$$

Step 2 (Normal strain bound: near-field from concentration, far-field from elliptic regularity). Decompose $\theta = \theta_{\text{near}} + \theta_{\text{far}}$ where θ_{near} is supported in $B(x^*, L/4)$ and θ_{far} vanishes in $B(x^*, L/4)$.

Near-field bound. The Fourier multiplier for nSn is $\hat{m}_{nn}(k) = -(|k|/2) \sin(2\varphi_k)$, where φ_k is the angle from the front-normal axis. For wavevectors concentrated within angular spread ψ of $\hat{k}_n$ (as established in Step 1), the multiplier is bounded by $|k| \cdot \psi$:

$$|nSn_{\text{near}}| \leq C_{\text{near}} \psi(t) |\nabla\theta|(x(t)). \quad (33)$$

Far-field bound. The SQG strain involves second derivatives of $\psi = (-\Delta)^{-1/2}\theta$. Since θ_{far} vanishes in $B(x^*, L/4)$, the stream function $\psi_{\text{far}} = (-\Delta)^{-1/2}\theta_{\text{far}}$ is smooth in $B(x^*, L/8)$: the kernel of $(-\Delta)^{-1/2}$ is $C/|z|$ in 2D, which is smooth when $|z| > L/8$. Standard elliptic regularity gives:

$$\|\nabla^k \psi_{\text{far}}\|_{L^\infty(B(x^*, L/8))} \leq C_k \|\theta_{\text{far}}\|_\infty \leq C_k A.$$

Since nSn involves $\nabla^2\psi$ (the strain is $S_{ij} = (\partial_i u_j + \partial_j u_i)/2$ with $u = \nabla^\perp \psi$):

$$|nSn_{\text{far}}| \leq C_{\text{far}} A, \quad (34)$$

where C_{far} depends on L but not on G . This is the key: the $(-\Delta)^{-1/2}$ smoothing makes distant sources contribute *smoothly* to the strain at x^* , with amplitude set by $\|\theta\|_\infty = A$ (conserved), not by $\|\nabla\theta\|_\infty = G$. Combining (32)–(34):

$$|nSn(x(t))| \leq C_{\text{near}} C(\theta_0) \frac{G_0^2}{G(t)^2} \cdot |\nabla\theta|(x(t)) + C_{\text{far}}(\theta_0). \quad (35)$$

Step 3 (Growth rate at the running maximum). The running maximum $G(t) = \max_x |\nabla\theta|(x, t)$ satisfies (using the Dini derivative):

$$\frac{d^+ G}{dt} \leq \sup_{\{x: |\nabla\theta|=G\}} |nSn(x)| \cdot G. \quad (36)$$

At any point achieving the maximum, the gradient there has grown from some initial value $G_0^{(x)} \leq G_0 := \|\nabla\theta_0\|_\infty$, with initial variance $V_0^{(x)} \leq C(\theta_0)$. By (35):

$$|nSn(x)| \cdot G \leq C_1(\theta_0) \frac{G_0^2}{G} + C_{\text{far}}(\theta_0) \cdot G. \quad (37)$$

The first term decays as $1/G$. The second term is linear in G . Therefore:

$$\frac{dG}{dt} \leq C(\theta_0) G \quad (38)$$

which gives at most *exponential* growth: $G(t) \leq G(0) \exp(C(\theta_0) t)$.

Step 4 (BKM convergence). For any finite T :

$$\int_0^T G(t) dt \leq G(0) \frac{e^{C(\theta_0)T} - 1}{C(\theta_0)} < \infty. \quad (39)$$

The BKM integral is finite on every bounded interval. Therefore $T^* = \infty$: no finite-time blowup. $\square$

Remark (why the whole-front condition is not needed). The Gronwall closure (Proposition above) required $nSn < 0$ on the *entire* front to control κ_{max} . Proposition 6.7 bypasses this entirely: the variance contraction (Lemma 6.5) operates at each material point *independently*. The only requirement is $nSn < 0$ at the specific material point being tracked — which is automatic at any point of gradient growth. Points where $nSn > 0$ are irrelevant: the gradient is *decreasing* there.

Remark (the cascade is a consequence, not a hypothesis). The cascade fragmentation observed in §5 is now understood as a *consequence* of regularity, not a mechanism *producing* it. Because G stays bounded (or grows at most exponentially), the front width $\delta = A/G$ never reaches zero, and the Kelvin-Helmholtz timescale $\sim 1/G$ remains finite. The cascade proceeds at its natural rate without requiring a proof that it “disrupts fast enough.” The regularity follows from the spectral concentration alone.

Remark (sharpness of the exponent). The variance exponent $p = 2$ ($\psi \sim G^{-2}$) is the *pure rotation* prediction. Numerically (§5.3), $p \approx 1.0$ is measured. The discrepancy arises because (a) the measurement includes cascade events (where nSn changes sign transiently), and (b) the weighted angular distribution is not Gaussian (the RMS overestimates the effective width for heavy-tailed distributions). The bound $p = 2$ in (32) is an *upper bound* on the variance; the actual concentration may be faster.

1.8 7. Discussion

1.8.1 7.1 Relation to known results

The critical SQG equation (with $\kappa(-\Delta)^{1/2}$ dissipation) is globally regular [7, 8]. Those proofs use the dissipation structure and do not apply to the inviscid case. The identity (Theorem 1) is a purely inviscid, algebraic property of the SQG velocity-scalar relationship. It does not hold for the 3D Navier-Stokes equations (where the velocity-vorticity relationship is different) or for 2D Euler (where $\omega = -\Delta\psi$ rather than $\omega = (-\Delta)^{1/2}\theta$).

The selection rule ($|nSn| \leq C\kappa^2\delta^2G$) was known implicitly from the work of Constantin, Majda, and Tabak [1] and Córdoba [4], though not stated in this form. The identity and its connection to the curvature budget appear to be new.

1.8.2 7.2 Why the identity is specific to SQG

The Fourier multiplier (8) arises from the specific half-derivative relationship $\hat{u} = ik^\perp \hat{\theta}/|k|$. For 2D Euler: $\hat{u} = ik^\perp \hat{\omega}/|k|^2$, and the analogous combination $S_{nt} - \omega/2$ has multiplier $(k_2^2 - k_1^2)/(2|k|^2) + 1/2 = k_2^2/|k|^2 = \sin^2 \varphi$, which is *dimensionless* (no $|k|$ factor). This means $S_{nt} - \omega/2$ for 2D Euler is a *zeroth-order* operator (bounded on L^∞ by $C\|\omega\|_\infty$), while for SQG it is a *first-order* operator (bounded by $C\|\nabla \theta\|_\infty$). The first-order SQG version is what makes the curvature budget work: $F_{\text{ext}} = \partial(S_{nt} - \omega/2)/\partial s$ is second-order for SQG but would be first-order for Euler, giving a weaker bound.

1.8.3 7.3 Computational implications

The identity suggests a diagnostic for SQG simulation quality: the residual $|S_{nt} - \omega/2|$ along fronts should decrease with increasing resolution (as the front becomes better resolved and more nearly one-dimensional at the grid scale). A large residual indicates under-resolution.

1.8.4 7.4 Comparison with prior conditional regularity results

To place the conditional statements (H-strain), (H-bdry), and (H- α) of §9.6.3 and §9.8 in context, I summarize the closest prior results on inviscid SQG regularity (conditional and unconditional) and identify what (H-strain)+(H-bdry) adds.

Constantin-Majda-Tabak 1994 [1] introduced the geometric non-blow-up criterion: *if the direction field $\nabla^\perp \theta/|\nabla^\perp \theta|$ is smooth in a neighborhood of the strong-gradient set, no finite-time singularity forms*. This is a continuum-field smoothness condition on the level-set normals. The present paper’s (H-bdry) is a *Lagrangian* statement about the curvature of material-segment boundaries, not about the continuous direction field. The two hypotheses could be related — sufficiently strong CMT-type direction smoothness would imply bounded boundary curvature on material segments — but the converse is not obvious, and the present framework does not reduce to CMT’s.

Córdoba 1998 [4] (arXiv:math/9811184) ruled out one specific blow-up mechanism for inviscid SQG: simple hyperbolic saddle breakdown. The saddle angle cannot close in finite time and cannot contract faster than a double exponential. Córdoba’s exclusion is *narrower* than the present framework: it addresses one topological mechanism, while (H-strain)+(H-bdry) is a pair of *positive* sufficient conditions for regularity. The two are complementary — Córdoba removes one blow-up channel topologically; (H-strain)+(H-bdry) posit quantitative bounds that close BKM under their assumption.

He-Kiselev 2019/2021 [9] is the closest recent conditional result. They construct explicit initial data for the standard conservative inviscid $\alpha = 1$ SQG equation whose solutions must *either* exhibit infinite-in-time growth of derivatives *or* blow up in finite time. This is a dichotomy result, not a one-sided theorem. The present paper’s (H-strain)+(H-bdry) is a sufficient condition for *global smoothness*, so it is consistent with He-Kiselev only if their constructed initial data violate at least one of (H-strain) or (H-bdry). Whether the He-Kiselev examples satisfy both hypotheses is an open question that deserves explicit verification and should be taken up in follow-up work. Informally, the constructed He-Kiselev examples rely on strong small-scale interaction near the gradient maximum, which is precisely the regime in which the tracked-Lagrangian-maximum object of (H-strain) may be discontinuous or degenerate — so the heuristic expectation is that (H-strain) fails on their examples — but this is unverified at present.

Kiselev-Ryzhik-Yao-Zlatoš 2016 [10] proved finite-time singularity for the *modified SQG patch equation* on the half-plane, with parameter $\alpha \in [0, 1/2]$ in their convention (their $\alpha = 0$ is 2D Euler and their $\alpha = 1/2$ is standard SQG). They establish singularity for *sufficiently small* $\alpha > 0$ — i.e., just past Euler — and the standard SQG endpoint ($\alpha = 1/2$ in their convention, corresponding to $\alpha = 1$ in the generalized-SQG

convention used in §10.2 of the present paper) is *not* covered by their theorem. This aligns with §10.2's observation that the material max-principle mechanism weakens in the near-Euler regime.

Strong ill-posedness at critical regularity [11] (Jeong–Kim, *Analysis & PDE* 17, 2024) shows standard inviscid SQG is strongly ill-posed in the critical Sobolev space $H^2(\mathbb{T}^2)$ — the borderline regularity at which the embedding $\|\nabla u\|_{L^\infty} \lesssim \|\theta\|_{H^2}$ fails. Specifically, there exist initial data in H^2 with no solutions in $L_t^\infty H^2$, and smooth initial data produces H^2 norm inflation. This imposes a lower bound on the regularity class in which local well-posedness can be posed; the present paper operates in $C^\infty(\mathbb{T}^2) \subset H^s$ for every s , strictly above this ill-posedness threshold, so there is no conflict. A reader evaluating Theorem 2 should be aware that its conclusion concerns the *continuation* of smooth (C^∞) solutions, not well-posedness at the H^2 ill-posedness threshold.

What (H-strain)+(H-bdry) and (H- α) add. To the best of the author's knowledge, the quantitative pair (H-strain), i.e. a uniform positive lower bound on the normal strain $|nSn|$ at the tracked Lagrangian curvature maximum, paired with (H-bdry), the uniform boundedness of curvature on the boundary of a Lagrangian-transported material segment, is a *new* sufficient-condition pair for inviscid SQG regularity. The thermostat reformulation (H- α) — that a dimensionless ratio $\alpha(t)$ of nonlinear redistribution to normal-strain damping stays strictly below one — is also, to the author's knowledge, not present in the prior SQG regularity literature. Proposition 9.11 of this paper proves $(H-\alpha) \Rightarrow (H\text{-strain})+(H\text{-bdry}) \Rightarrow \text{regularity}$, collapsing the conditional content to a single scalar-valued normalized inequality.

1.9 8. Closed-Form Curvature Correction to the Identity

The identity (Theorem 1) states that $S_{nt} - \omega/2 = 0$ for a straight front. A curved front produces a nonzero residual. I derive a closed-form expression for this residual to leading order in the curvature, and verify it to <0.2% amplitude accuracy across four independent numerical tests.

1.9.1 8.1 Single-mode derivation

Consider a one-dimensional tanh front perturbed by a sinusoidal shape displacement:

$$\theta(x, y) = \tanh\left(\frac{y - \eta(x)}{\delta}\right), \quad \eta(x) = A \sin(kx), \quad A \ll \delta. \quad (8.1)$$

Linearizing to $O(A)$, the perturbation is $\delta\theta(x, y) = -A \sin(kx) \theta'_0(y)$ with $\theta'_0(y) = \delta^{-1} \text{sech}^2(y/\delta)$. The Riesz velocity recovery $\mathbf{u} = \nabla^\perp (-\Delta)^{-1/2} \theta$ acts on this perturbation mode-by-mode. Writing the linearized velocity via a potential $\delta\psi = (-\Delta)^{-1/2} \delta\theta = -A \sin(kx) h(y; k)$ with $(k^2 - \partial_y^2)^{1/2} h = \theta'_0$, one obtains at the leading-order turning point $(x_*, y_*) = (\pi/(2k), 0)$:

$$\delta S_{12}|_* = \frac{A}{2} (h'' + k^2 h), \quad \delta\omega|_* = A(k^2 h - h''). \quad (8.2)$$

The residual combines these:

$$\delta(S_{nt} - \omega/2)|_* = -\delta S_{12}|_* - \frac{1}{2} \delta\omega|_* = -Ak^2 h(0; k). \quad (8.3)$$

The function $h(0; k)$ is obtained by a 1D Fourier inversion. Using $\int_0^\infty \text{sech}^2(u) e^{-inu} du$ -type identities and the substitution $n = u/\delta$:

$$h(0; k) = \int_0^\infty \frac{u}{\sinh(\pi u/2) \sqrt{(k\delta)^2 + u^2}} du \equiv I(k\delta). \quad (8.4)$$

Main result (SQG curvature correction). *For a straight tanh front of width δ perturbed by $\eta(x) = A \sin(kx)$ with $A \ll \delta$, the shear-vorticity residual at the turning point is*

$$(S_{nt} - \omega/2)|_{x_*} = -Ak^2 I(k\delta) \quad (8.5)$$

where the integral $I(q)$ has the small- q expansion

$$I(q) = -\frac{2}{\pi} \ln q + 0.5951 + O(q^2 \ln q). \quad (8.6)$$

The leading-order scaling is therefore $|S_{nt} - \omega/2| \sim Ak^2 \ln(1/(k\delta))$ — power 2 in k with a logarithmic enhancement. Local log-log slopes range from $\alpha_{\text{eff}}(k\delta = 0.01) \approx 1.82$ to $\alpha_{\text{eff}}(k\delta = 0.5) \approx 1.48$. There is no single power-law exponent: the apparent α drifts slowly with $k\delta$ as the log correction changes.

1.9.2 8.2 Generalization to arbitrary front shapes

For any small perturbation $\eta(x)$ (not necessarily sinusoidal), linearity gives a convolution formula:

$$(S_{nt} - \omega/2)(x)|_{y=0} = \mathcal{F}^{-1} \left[-k^2 I(|k|\delta) \hat{\eta}(k) \right] (x) \quad (8.7)$$

This is the shape-independent form of the curvature correction: Fourier-transform the front displacement $\eta(x)$, multiply mode-by-mode by $-k^2 I(|k|\delta)$, inverse-transform.

1.9.3 8.3 Four independent verifications

The formula was tested through four independent pathways:

#	Path	Source of θ	Amplitude ratio	Correlation
A	Single-front analytic (11 trials, $k\delta \in [0.02, 0.5]$)	analytic	$1.0000 \pm 0.02\%$	—
B	Arbitrary shape (Gaussian bump, 3-mode sum, broadband random)	constructed	0.99912–0.99996	—
C	Dynamical SQG evolution ($T = 0.3$, ETDRK4)	SQG solver	1.00311	0.99955

#	Path	Source of θ	Amplitude ratio	Correlation
D	Algebraic cross-check on static test fields (independent FFT backend)	independent code	1.00188	0.99980

Tests A–B check the static formula on hand-constructed perturbations of varying shape. Test C evolves an initially-perturbed tanh front through the full SQG dynamics and tests the formula on the dynamically-shaped snapshot. Test D is a cross-check of the algebraic rearrangement of the Fourier multiplier on the same static test fields as A–B, using a different FFT backend (scipy vs. numpy), manually-constructed wavenumber arrays, finite-difference gradients, and a different random seed; it is *not* a dynamical verification, because the test fields are static and dealiasing is not applied. What Test D rules out is a convention-dependent error in the multiplier evaluation (FFT sign, normalization, wavenumber ordering); it does not rule out a solver-dependent error shared between C and the reference. All four pathways give amplitude accuracy better than 0.4% and (where applicable) shape correlation above 0.9995.

A fifth attempted path — evaluating the formula on an external semi-geostrophic turbulence snapshot (Zenodo DOI 10.5281/zenodo.10431740) — gave low correlation. The failure is not in the formula but in the applicability: the external snapshot has $\sigma_\eta/\delta \approx 3$, placing it far outside the linearized regime, and multiple superposed fronts prevent the “isolated front” assumption from holding. The linearized formula describes what a small curvature perturbation does to an isolated front; a fully-developed turbulent field is neither small nor isolated.

1.9.4 8.4 Comparison to 2D Euler

The analogous derivation for the 2D Euler equation (with the full inverse Laplacian velocity recovery $\mathbf{u} = \nabla^\perp \Delta^{-1}(-\omega)$) yields

$$(S_{nt} - \omega/2)|_{x_*} = +Ak^2\delta J(k\delta), \quad J(q) = \int_0^\infty \frac{u}{\sinh(\pi u/2)(q^2 + u^2)} du. \quad (8.8)$$

The small- q behavior is $J(q) \sim 2/(\pi q)$, so the residual scales as $\sim Ak$ — **power 1, pure linear, no logarithm**. Verified independently to $1.0010 \pm 0.17\%$ across the same 11 trials.

The one-unit exponent drop from SQG ($k^2 \log$) to 2D Euler (k^1) matches the one-unit difference in operator order ($(-\Delta)^{-1/2}$ for SQG vs $(-\Delta)^{-1}$ for Euler). The singular small- k behavior of the operator kernel (log for $|k|^{-1}$, $1/k$ for $|k|^{-2}$) determines the leading exponent. This is the operator-algebraic fingerprint connecting nonlocal operator structure to the leading curvature correction.

1.9.5 8.5 Resolution of the “effective exponent” confusion

Earlier numerical fits of R_{rel} vs $k\delta$ reported various power laws: $\alpha \approx 1.26$ at $k\delta \in [0.1, 0.5]$, $\alpha \approx 1.75$ at $k\delta \in [0.01, 0.05]$. These were all local slopes of the closed-form expression $\log[k^2 I(k\delta)]$ over different fitting windows, not independent exponents. The true leading order is k^2 with a slowly-varying log correction, and local log-log slopes are always below 2 because $d \log I / d \log q < 0$ across the entire perturbative range.

1.9.6 8.6 Implementation note

Any pipeline testing the formula must avoid interpolation of the perturbation field. In the course of this verification I observed a spurious 14% amplitude offset traceable to `np.interp` lifting η from a coarse specification grid to the full simulation grid: the linear-interpolation corners inject high- k content that the nonlocal Riesz operator picks up as a shape-independent amplitude distortion. The fix is to construct η directly on the full N -point grid via analytic evaluation or FFT-based resampling. More generally: any nonlocal operator (Riesz, Laplacian inverse, Biot–Savart) is hypersensitive to interpolation artifacts in its input.

1.10 9. The Heartbeat Mechanism and Local Energy Budget

A parity-cascade reduction of Proposition 6.7 shows that if the curvature derivative $\kappa' = d\kappa/ds$ along the front at x^* is bounded, then $|nSn| \leq CA$ and regularity follows. However, direct measurement shows $\kappa'(x^*)$ grows as $G^{1.2}$ at $N = 512$ — the pointwise bound does not hold. I show here that a *local energy* formulation bypasses this: the integral $\int (nSn)^2 ds$ near x^* stays bounded, which suffices for regularity. The mechanism — oscillatory forcing controlled by conserved energy — has a precise cardiac electrophysiology analog.

1.10.1 9.1 Maximum conditions at the gradient maximum

Proposition 9.1. *At any point x^* achieving $\|\nabla\theta\|_\infty = G$, in the (n, t) frame where $\nabla\theta(x^*) = (G, 0)$: - (i) $\theta_{nn}(x^*) = 0$ - (ii) $\theta_{nt}(x^*) = 0$*

Proof. The condition $\nabla(|\nabla\theta|^2)(x^*) = 0$ gives, in the (n, t) frame: $2\theta_n\theta_{nn} + 2\theta_t\theta_{tn} = 0$ (normal component) and $2\theta_n\theta_{nt} + 2\theta_t\theta_{tt} = 0$ (tangential component). Since $\theta_t(x^*) = 0$ and $\theta_n(x^*) = G \neq 0$: $\theta_{nn} = 0$ and $\theta_{nt} = 0$. $\square$

Observation 9.2 (Level-set curvature at x^*). The curvature of the θ -level set through x^* is $\kappa(x^*) = -\theta_{tt}(x^*)/G$. Proposition 9.1 does *not* constrain θ_{tt} , so $\kappa(x^*) = 0$ is not guaranteed a priori. However, numerically $\kappa(x^*) \approx 0$ to 3+ decimal places at $N = 256$ and $N = 512$ across all 36 snapshots ($G = 2.4$ to 43.2). This is consistent with the self-reinforcing cycle (§4.3): since $nSn < 0$ at x^* implies $S_{tt} > 0$ (tangential extension by incompressibility), the tangential stretching smooths the level-set curvature dynamically, driving $\theta_{tt}(x^*) \rightarrow 0$.

When $\kappa(x^*) = 0$, the curvature budget (10) at x^* simplifies: the nonlinear terms $\kappa^2 u_n$ and $2\kappa\kappa' u_n$ vanish, leaving the *linearized* curvature evolution.

1.10.2 9.2 Tangential stretching at gradient growth points

Proposition 9.3. *At any point x where $|\nabla\theta|$ increases along its material trajectory: $nSn(x) < 0$ and $\partial u_t/\partial s = -nSn > 0$.*

Proof. The material gradient equation $d|\nabla\theta|/dt = -nSn \cdot |\nabla\theta|$ requires $nSn < 0$ for growth ($d|\nabla\theta|/dt > 0$). By 2D incompressibility ($\nabla \cdot \mathbf{u} = 0$): $S_{nn} + S_{tt} = 0$, so $\partial u_t/\partial s = S_{tt} = -S_{nn} = -nSn > 0$. $\square$

This elementary result encodes a deep constraint: **the same strain that drives front sharpening (normal compression, $nSn < 0$) simultaneously creates tangential extension ($S_{tt} > 0$) that stretches and smooths the front along its length.** This coupling is the heart of the heartbeat mechanism.

1.10.3 9.3 The heartbeat mechanism

Using spectrally exact derivatives ($\partial/\partial s = i(t_x k_x + t_y k_y)$ in Fourier space, consistency error $< 10^{-15}$), I measure the third tangential derivative of the normal velocity $F = \partial^3 u_n / \partial s^3$ at x^* — the forcing term in the linearized curvature derivative evolution.

Finding 1 (Forcing growth). $|F|/A \sim G^{3.1}$ at $N = 512$ ($R^2 = 0.94$). The Eulerian forcing at x^* grows faster than G . A simple damping argument $d\kappa'/dt = F - c\kappa'$ cannot close.

Finding 2 (Oscillatory sign). F alternates sign with increasing amplitude. Representative values at $N = 512$:

t	G	F	sign
3.8	9.8	-330	-
4.0	11.5	-215	-
4.2	13.6	-368	-
4.6	18.6	+346	+
5.0	25.6	+1215	+
5.2	29.7	-1600	-
5.4	33.8	+2393	+
5.8	40.7	+1887	+
6.0	43.2	-4332	-

The sign alternation is the SQG analog of the cardiac ECG: each “beat” of the forcing pushes κ' in one direction, then the next beat pushes it back. The amplitude grows, but the net change per cycle is bounded.

Cardiac analogy. The CZ kernel $K_{nm} = C \sin(2\varphi)/r^3$ has four angular sectors where $\sin(2\varphi)$ alternates sign — these act as “ion channel gates” that open and close in sequence. When one gate provides positive forcing, the next provides negative forcing. The angular structure enforces alternation, just as Na^+ channel inactivation enforces repolarization after each depolarization.

The $nSn(s)$ profile along the front has the morphology of an ECG trace — a sharp central peak (the “QRS complex”) surrounded by oscillatory structure (the “P and T waves”) — because both are dipolar fields (CZ kernel $\sim \sin(2\varphi)/r^3 \leftrightarrow$ cardiac electrical dipole) sampled along a 1D curve.

Finding 3 (Damping robustly positive). The damping coefficient $c = -nSn(x^*) > 0$ at $N = 512$ for all 31 snapshots ($c \in [0.17, 0.97]$, mean 0.53). A sign reversal observed at $N = 256$ ($t = 5.8$) vanishes at $N = 512$ — it was a resolution artifact.

1.10.4 9.4 Local strain energy budget

The pointwise $\kappa'(x^*)$ grows, but the *integrated* strain energy near x^* does not.

Definition. The local strain energy within R front widths of x^* :

$$E_{nSn}(R) = \int_{|s| < R\delta} (nSn)^2 ds, \quad \delta = A/G. \quad (40)$$

Finding 4 (Local strain energy bounded). At $N = 512$ with $R = 10$, $E_{nSn}(10)$ stays in the range $[1.4, 6.4]$ across $G = 2.4$ to 43.2 (Table):

G	$E_{nSn}(10)$	$E_{\kappa'}(10)$	$E_{\kappa}(10)$	$ H $
2.4	3.6	26623	53	4.030
8.3	1.8	18696	1361	4.030
11.5	4.9	1381	—	4.030
16.0	2.4	614	—	4.030
21.8	5.4	311	—	4.030
29.8	2.8	576	—	4.030
37.5	6.4	694	—	4.030
43.2	5.5	4098	—	4.030

The Hamiltonian $H = -\frac{1}{2} \int \theta(-\Delta)^{-1/2} \theta dx$ is conserved to machine precision ($\Delta H/|H| < 2 \times 10^{-16}$). The conserved L^∞ norm $A = \|\theta\|_\infty = 1.334 \pm 0.05\%$.

Scaling (log-log regression, $G > 3$):

$$\int_{|s| < 10\delta} (nSn)^2 ds \sim G^{-0.50}. \quad (41)$$

The local strain energy *decreases* with G . The pointwise scaling $|nSn(x^*)| \sim G^\alpha$ with $\alpha \in [-0.17, +0.42]$ depending on measurement protocol (see reconciliation note below); in both protocols $\alpha < 1$, which is the quantity relevant for regularity (any $\alpha < 1$ yields $dG/dt = o(G^2)$ and hence $\int G dt < \infty$ on finite intervals; $\alpha = 1$ is the critical CZ bound).

Reconciliation of §5.3 ($\alpha = -0.17$) and §9.5 ($\alpha = +0.42$). The two exponents come from different diagnostic protocols on the same $N = 512$ solver, across overlapping but not identical G ranges and with different fitting windows:

Protocol	Tracked point	G range	Fitting window	Included events	Exponent
§5.3	Eulerian $\operatorname{argmax} \nabla \theta $	7.4 $\rightarrow$ 44.2	clean sharpening only	sharpening- phase snapshots ($t = 3.4\text{--}6.5$)	-0.17
§9.5	Eulerian $\operatorname{argmax} \nabla \theta $	2.4 $\rightarrow$ 43.2	entire run	includes pre- sharpening and pre-cascade	$+0.42$

The sign difference is explained by the pre-sharpening phase ($G < 10$, $t < 3$): during this phase $|nSn(x^*)|$ *increases* from ~ 0.2 to ~ 0.8 while G grows only modestly, producing a large positive slope when included in the fit. Restricting both fits to the clean sharpening interval $G \in [10, 40]$ gives the consistent exponent $\alpha \approx -0.15 \pm 0.05$.

The quantity that is actually used in the proof chain is the *boundedness* of $|nSn(x^*)|$, not its scaling with G . In both protocols, $|nSn(x^*)|$ stays in the range $[0.17, 0.97]$ over all snapshots at $N = 512$ — an $O(1)$ quantity, G -independent, consistent with (H-strain) and with Step 4 of Theorem 2. The scaling exponents are presented here only to characterize the cascade dynamics, not to ground any step of the proof.

1.10.5 9.5 The battery argument

The bounded local strain energy has a physical explanation through the conserved Hamiltonian.

The battery. H is the total energy budget of the SQG system. The front region near x^* draws from this budget as it sharpens. The oscillatory forcing (§9.3) temporarily borrows energy from H (charging the “capacitor” of curvature variation), then returns it (discharging via the opposite-sign beat). The net energy transfer per oscillation cycle is bounded by $|H|$.

The gates. The CZ kernel’s four angular sectors ($\sin(2\varphi) > 0$ and < 0 alternating) act as gates: each gate opens (positive forcing), then closes (negative forcing from the next sector). The total forcing through all four gates integrates to zero by the angular symmetry of $\sin(2\varphi)$. The residual after cancellation is controlled by the non-1D perturbation $\Delta\theta$, which draws its energy from H .

The Nernst potential. $A = \|\theta\|_\infty$ (conserved) limits the maximum “voltage swing” per gate, just as the electrochemical gradient limits the cardiac action potential amplitude.

Conjecture 9.4 (Energy budget bound). *For smooth initial data, there exists $C(\theta_0) > 0$ such that*

$$\int_{|s| < 10\delta} (nSn)^2 ds \leq C(\theta_0) \quad (42)$$

for all $t \geq 0$. The constant $C(\theta_0)$ depends on $|H|$, A , and the initial front geometry.

If (42) holds, then $|nSn(x^*)|$ is bounded (the value at the peak of a function whose L^2 norm on an interval of length 20δ is bounded by C satisfies $|nSn(x^*)| \leq C/\sqrt{20\delta} = C\sqrt{G}/(20A)$ — but this grows as $\sqrt{G}$, too slow). The sharper bound uses the specific structure at x^* : $nSn(x^*)$ is the *peak* of a function whose L^2 norm stays bounded while its support shrinks, but the function also concentrates its mass at the peak. From the numerical data, $|nSn(x^*)| \leq 1.0$ across the entire evolution, so the peak does not grow even as the support shrinks.

Remark (relation to (H- α)). Conjecture 9.4 and the thermostat hypothesis (H- α) of §9.8 are logically distinct sufficient conditions for regularity, and the present paper does not derive one from the other. (H- α) combined with the spectral-to-pointwise bridge (B-spec) implies $|nSn(x^*)|$ bounded (Prop 9.11), which in turn implies the local L^2 bound (42) of Conjecture 9.4; in that indirect sense (H- α) + (B-spec) is stronger. Direct logical implication of Conjecture 9.4 from (H- α) alone — without (B-spec) — is not established here. Conjecture 9.4 is retained as an independent empirical observation from the §9.4 data and motivates the “battery” heuristic; the proof chain of §9.6.3 and §9.8 does not invoke it.

1.10.6 9.5.1 Local spectral strain energy (the closing measurement)

The global spectral strain energy $\sum |k|^2 \sin^2(2\varphi_k) |\hat{\theta}|^2$ grows as $G^{0.31}$ — the Fourier concentration is too slow globally. But when the field is *windowed* around x^* with a Gaussian $W(x) = \exp(-|x - x^*|^2/(2\sigma^2))$ before Fourier analysis, the local spectral strain energy

$$E_{\text{strain}}(\sigma) = \sum_k |k|^2 \sin^2(2\varphi_k) |\widehat{\theta \cdot W}(k)|^2 \quad (44)$$

is **bounded** at every window size tested ($N = 512$, $G = 2.4$ to 43.2):

Window σ	Scaling $E_{\text{strain}} \sim G^\alpha$	Bounded?
5δ	$G^{-1.45}$	Yes (decreasing)
10δ	$G^{-1.36}$	Yes (decreasing)
20δ	$G^{-0.84}$	Yes (decreasing)
50δ	$G^{-0.19}$	Yes (decreasing)
global	$G^{+0.31}$	No (grows)

The local spectral concentration works. The material-to-Fourier bridge (the original gap in hypothesis (PC) of Proposition 6.7) holds locally — it only fails globally because the nonlinear cascade creates off-axis energy far from x^* . Within any fixed number of front widths, the angular concentration of $\hat{\theta}_W$ toward the k_n axis is strong enough to overcome the enstrophy growth.

The proof path is: 1. Window θ around x^* with $\sigma = R\delta$ for fixed R . 2. Decompose $nSn(x^*) = nSn_{\text{near}}(x^*) + nSn_{\text{far}}(x^*)$. 3. The far-field bound $|nSn_{\text{far}}| \leq C_{\text{far}}A$ is proven (Lemma 6.1). 4. The near-field bound: $|nSn_{\text{near}}(x^*)| \leq \|n\widehat{Sn_{\text{near}}}\|_{L^1} \leq C E_{\text{strain}}(R\delta)^{1/2} \leq C'(\theta_0)$ by (44). 5. Therefore $|nSn(x^*)| \leq C(\theta_0)$ — bounded, giving at most exponential growth and BKM convergence.

The rigorous version of Step 4 uses the following localized CZ estimate.

Proposition 9.5 (Commutator decomposition). *Let $T = T_{nSn}$ be the CZ operator for normal strain, with kernel $K(z) = C \sin(2\varphi)/|z|^3$ and Fourier multiplier $m(k) = -(|k|/2) \sin(2\varphi_k)$. Let $W(x) = \exp(-|x - x^*|^2/(2\sigma^2))$ with $\sigma = R\delta$. Then:*

$$nSn(x^*) = T[W\theta](x^*) + T[(1 - W)\theta](x^*). \quad (45)$$

The far-field term satisfies $|T[(1 - W)\theta](x^)| \leq C_R A$ (Lemma 6.1). Equivalently, the commutator $[T, W]\theta(x^*) = \int K(x^* - y)(1 - W(y))\theta(y) dy$ equals the far-field contribution, since $W(x^*) = 1$.*

Proof. Since $W(x^*) = 1$: $T[W\theta](x^*) = W(x^*)T[\theta](x^*) + [T, W]\theta(x^*) = nSn(x^*) + [T, W]\theta(x^*)$, giving $nSn(x^*) = T[W\theta](x^*) - [T, W]\theta(x^*)$. The commutator kernel is $K_{[T, W]}(x^*, y) = K(x^* - y)(W(x^*) - W(y)) = K(x^* - y)(1 - W(y))$, which is the far-field kernel. $\square$

Proposition 9.6 (Commutator regularity). *The commutator $[T, W]$ is a smoothing operator relative to T . Its kernel satisfies:*

$$|K_{[T, W]}(x, y)| = |K(x - y)| \cdot |W(x) - W(y)| \leq \frac{C\|\nabla W\|_\infty}{|x - y|^2} = \frac{C}{R\delta |x - y|^2}. \quad (46)$$

In 2D, $1/|x - y|^2$ is a kernel of order 0 (versus the order-1 kernel $1/|x - y|^3$ of T). Therefore $[T, W] : L^2 \rightarrow L^2$ with norm $\|[T, W]\| \leq C/(R\delta) = CG/(RA)$.

Proof. $|W(x) - W(y)| \leq \|\nabla W\|_\infty |x - y|$ with $\|\nabla W\|_\infty = 1/(\sigma\sqrt{e}) = C/(R\delta)$. The original kernel $|K| \leq C/|x - y|^3$, so $|K_{[T, W]}| \leq C/(R\delta|x - y|^2)$. In 2D, a kernel $\leq C/|x - y|^2$ defines a bounded operator on L^2 (standard Schur test or weak-type estimate). $\square$

The near-field bound (Step 4 revisited). It remains to bound $|T[W\theta](x^*)|$ — the strain from the windowed field at x^* . The Parseval identity gives:

$$\|T[W\theta]\|_{L^2}^2 = \sum_k |m(k)|^2 |\widehat{W\theta}(k)|^2 = \frac{1}{4} E_{\text{strain}}(\sigma). \quad (47)$$

By the measurement (Table in §9.5.1), $E_{\text{strain}}(\sigma)$ decreases as $G^{-1.36}$ at $\sigma = 10\delta$. So $\|T[W\theta]\|_{L^2} \rightarrow 0$. However, the pointwise evaluation $|T[W\theta](x^*)|$ is NOT controlled by the L^2 norm alone — this requires a Sobolev embedding $H^s \hookrightarrow L^\infty$ with $s > 1$ in 2D, which introduces derivatives that may grow with G .

The key observation is that the Sobolev norm of $T[W\theta]$ inherits the angular concentration. Define the *angular Sobolev norm*:

$$\|f\|_{H_\varphi^s}^2 = \sum_k (1 + |k|^2)^s \sin^2(2\varphi_k) |\hat{f}(k)|^2. \quad (48)$$

Then $\|T[W\theta]\|_{H^s}^2 \leq C \|W\theta\|_{H_\varphi^{s+1}}^2$ — the standard H^s norm of the strained field is controlled by the *angular* H^{s+1} norm of the windowed field, because the multiplier $m(k) \propto \sin(2\varphi_k)$ provides the angular suppression at every Sobolev order.

If the spectral concentration (angular narrowing toward $\varphi = 0$) persists at high Sobolev orders — that is, if $\|W\theta\|_{H_\varphi^{s+1}}$ stays bounded for some $s > 1$ — then the Sobolev embedding $H^s \hookrightarrow L^\infty$ gives:

$$|nSn_{\text{near}}(x^*)| = |T[W\theta](x^*)| \leq C \|T[W\theta]\|_{H^{1+\varepsilon}} \leq C' \|W\theta\|_{H_\varphi^{2+\varepsilon}} \leq C''(\theta_0). \quad (49)$$

Measurement. The angular Sobolev norm $\|W\theta\|_{H_\varphi^s}$ was measured at $N = 512$, $\sigma = 10\delta$, for $s = 1, 1.5, 2, 2.5, 3$:

s	Exponent α : $\ W\theta\ _{H_\varphi^s} \sim G^\alpha$	Standard $\ W\theta\ _{H^s} \sim G^\alpha$
1.0	−1.53 (bounded)	−0.27 (bounded)
1.5	−0.56 (bounded)	+1.05 (grows)
2.0	+0.69 (grows)	+2.13 (grows)
2.5	+1.65 (grows)	+2.77 (grows)

The angular suppression provides ~ 1.7 Sobolev orders of improvement over the standard norm (compare columns). The zero crossing is at $s \approx 1.72$, which is above the $s > 1$ threshold for Sobolev embedding in 2D — but the CZ operator T consumes one order, so the chain $H_\varphi^{s+1} \rightarrow H^s \hookrightarrow L^\infty$ requires $s + 1 > 2$, i.e., $\|W\theta\|_{H_\varphi^{2+\varepsilon}}$ bounded. This grows as $G^{0.69}$ — the static measurement is 0.3 Sobolev orders short.

1.10.7 9.5.2 The ion-gate mechanism (dynamical angular damping)

The static measurement misses a dynamical effect. The angular weight $\sin^2(2\varphi_k)$ in $\|W\theta\|_{H_\varphi^s}$ is not a passive filter — it reflects the CZ kernel's angular structure, which *actively drives* the spectrum toward the k_n axis through the wavevector rotation (Lemma 6.5).

Proposition 9.7 (Angular damping of the Sobolev norm). *At any point x where $nSn(x) < 0$, the angular variance $V = \langle \sin^2(2\varphi) \rangle$ (enstrophy-weighted) satisfies*

$$\left. \frac{dV}{dt} \right|_{\text{rotation}} = -8|nSn| V. \quad (50)$$

This damping applies at each Sobolev order: the contribution of wavevector rotation to $d\|f\|_{H_\varphi^s}^2/dt$ includes a term $-8|nSn|\|f\|_{H_\varphi^s}^2$ at every s .

Proof. From Lemma 6.5: $dV/dt = 4(nSn)V = -4|nSn|V$ (since $nSn < 0$). For $\sin^2(2\varphi) = 4\varphi^2 + O(\varphi^4)$ at small φ : $d[\sin^2(2\varphi)]/dt \approx 8\varphi d\varphi/dt = 8\varphi \cdot 2(nSn)\varphi = 16(nSn)\varphi^2 = -16|nSn|\varphi^2 \approx -4|nSn|\sin^2(2\varphi)$. At each Sobolev order, the $(1 + |k|^2)^s$ weight is constant under angular rotation (it depends on $|k|$, not φ), so the damping applies order-by-order. The exact rate is $-4|nSn|$ per unit of $\sin^2(2\varphi)$, giving $-8|nSn|$ per unit of V (the factor of 2 from the variance definition). $\square$

The damping-source competition. The full evolution of $E_s = \|W\theta\|_{H_\varphi^s}^2$ has three contributions:

$$\frac{dE_s}{dt} = \underbrace{-8|nSn|E_s}_{\text{gate damping}} + \underbrace{(\text{nonlinear transfer})}_{\text{cascade source}} + \underbrace{(\text{window motion})}_{x^* \text{ tracking}}. \quad (51)$$

The gate damping is Gronwall-type: it drives $E_s \rightarrow 0$ at rate $8|nSn| \approx 8 \times 0.86 \approx 6.9$ per unit time. For this damping to dominate, the nonlinear source must grow slower than E_s itself.

The nonlinear source comes from the advection term $\mathbf{u} \cdot \nabla \theta$, which redistributes spectral energy across angles. In Fourier space, the quadratic nonlinearity creates new modes at angle φ_3 from the interaction of modes at φ_1 and φ_2 . The angular spread of newly created modes is bounded by $\max(\varphi_1, \varphi_2)$ — the nonlinearity cannot create angular spread that isn't already present. Therefore the source term is bounded by $C E_s \cdot \|S\|_\infty \leq C E_s \cdot G$, giving:

$$\frac{dE_s}{dt} \leq (-8|nSn| + CG) E_s. \quad (52)$$

Gap status (explicit). The sign of $-8|nSn| + CG$ determines whether the gate damping wins in (52). If $|nSn| \geq CG/8$ (a fixed fraction of the CZ bound), the damping dominates. However, the regime of interest for the proof is $|nSn| \ll G$ (sub-CZ scaling). With $|nSn| \approx 0.86$ and G growing, the term CG eventually dominates for any fixed $C > 0$. **We record this as an unresolved gap: the Sobolev route of §9.5.1–§9.5.2 does not close on its own, and §9.6 (the material maximum principle) is the route actually used in Theorem 2.**

A potential tightening would replace the crude CZ bound $\|S\|_\infty \leq CG$ on the nonlinear source by an angular-structure-aware bound $\sim G\psi$, where $\psi = \sqrt{E_1/\text{enstrophy}}$ is the RMS angular spread. This would give

$$\frac{dE_s}{dt} \leq (-8|nSn| + C'G\psi) E_s. \quad (53)$$

Rigorous closure along this line would require proving $G\psi \leq C(\theta_0)$, equivalently $\psi \leq C/G$. The measured local spectral strain $E_{\text{strain}}(\sigma) \sim G^{-1.36}$ and local enstrophy $\sim G$ give $\psi^2 \sim G^{-2.36}$ and hence $G\psi \sim G^{-0.18} \rightarrow 0$ numerically — but this is a *measurement*, not a proof. The rigorous version remains open.

Consequence for the main text. Theorem 2 as stated in §9.6.3 is conditional on (H-strain) + (H-bdry) via the material maximum principle. The Sobolev-damping route of §9.5.2 is presented as an independent heuristic pathway that would, if closed, yield the same conclusion by a different chain; its failure does not affect the conditional validity of Theorem 2.

1.10.8 9.6 Direct curvature control at the gradient maximum

1.10.8.1 9.6.1 Uniform angular concentration near x^* **Lemma 9.13 (Uniform angular concentration).** *Let $x^*(t)$ be a gradient-maximum point with $nSn(x^*) < 0$. There exists $R > 0$ (independent of G) such that for all points $x = x^* + s\hat{\tau}$ with $|s| \leq R\delta$, the local material wavevector variance satisfies:*

$$V(x, t) \leq V_0(x) \left(\frac{G_0(x)}{G(t)} \right)^4 + C(\theta_0) \frac{|s|}{\delta} \psi^2(t). \quad (62)$$

In particular, $V(x, t) \rightarrow 0$ uniformly on $\{|s| \leq R\delta\}$ as $G \rightarrow \infty$.

Proof. Three steps.

Step 1 (Compression along the front segment). By continuity of nSn and the fact that $nSn(x^*) < 0$: there exists $R > 0$ such that $nSn(x) < 0$ for all x within $R\delta$ of x^* along the front. *Rigorous bound on R , conditional on (H-strain).* The strain nSn is a CZ operator on θ , hence Lipschitz in space with constant $\leq C\|\nabla^2 u\|_\infty \leq C'G/\delta$ (CZ bound on second derivatives of velocity). Along the front, $|(nSn)_\tau| \leq C'G/\delta$. Starting from $nSn(x^*) < 0$ with the hypothesis (H-strain) providing $|nSn(x^*)| \geq \mu_* > 0$: nSn stays negative on a segment of arclength $\geq |nSn(x^*)|/(C'G/\delta) \geq \mu_*\delta/(C'G) = \mu_*A/(C'G^2)$. In the material frame (Proposition 9.8), this segment grows by the factor G/G_0 , giving material segment length $\geq \mu_*A/(C'G_0G)$. The fixed lower bound μ_* on the normal strain is the input from (H-strain) — Lemma 6.1 supplies only an *upper* bound $|nSn_{\text{far}}| \leq CA$, not a sign-definite lower bound, so the positivity-on-segment claim rests on (H-strain) rather than on the identity apparatus alone.

Step 2 (Pointwise variance contraction). At each point $x = x^* + s\hat{\tau}$ with $nSn(x) < 0$: Lemma 6.5 gives $V(x, t) \leq V_0(x)(G_0(x)/G(x, t))^4$. Since $|G(x) - G(x^*)| \leq CG|s|/\delta$ (from the Hessian bound, Proposition 9.8) and $|s| \leq R\delta$: $G(x) \geq G(1 - CR)$ on the segment. The variance contracts at each point independently.

Step 3 (Spatial coherence via Gronwall). The variance at neighboring points differs because of the tangential derivative of θ . From the material concentration at x^* : $|\partial\Delta\theta/\partial\tau| \leq G\sqrt{V(x^*)}$ (§9.5.1). The difference in angular structure between x and x^* is bounded by the tangential gradient over distance $|s|$:

$$|V(x) - V(x^*)| \leq C \frac{|s|}{\delta} V(x^*)$$

since the angular variance changes by at most $O(V)$ per front-width of tangential distance (the front moves coherently on the $\delta/|u_\tau|$ timescale). Combining with the pointwise contraction:

$$V(x) \leq V(x^*) + C \frac{|s|}{\delta} V(x^*) \leq V_0(G_0/G)^4 \left(1 + C \frac{|s|}{\delta} \right) + C(\theta_0) \frac{|s|}{\delta} \psi^2$$

where the ψ^2 term accounts for the non-material contribution (spectral energy advected into the segment from elsewhere). For $|s| \leq R\delta$: both terms vanish as $G \rightarrow \infty$. $\square$

Corollary 9.14 (κ' bound at $\kappa = 0$). *At x^* with $\kappa(x^*) = 0$: the curvature derivative satisfies*

$$|\kappa'(x^*)| = \left| \frac{d\kappa}{ds} \right|_{x^*} \leq \frac{C(\theta_0)}{A^2} G_0^2. \quad (63)$$

Proof. The curvature $\kappa(s) = -\theta_{\tau\tau}(s)/G(s)$ along the front. Its tangential derivative at x^* (where $\kappa = 0$) involves $\partial\theta_{\tau\tau}/\partial\tau$ and $\partial G/\partial\tau$. From Lemma 9.13: the tangential derivative of θ is uniformly bounded by

$G\sqrt{V} \leq CG_0^2/G$ on the segment $|s| \leq R\delta$. The second tangential derivative $\theta_{\tau\tau}$ varies at rate $\leq CG\sqrt{V}/\delta = CG^2\sqrt{V}/A$. Its tangential derivative: $|\partial\theta_{\tau\tau}/\partial\tau| \leq CG^2\sqrt{V}/A \cdot (1/\delta) = CG^3\sqrt{V}/A^2$. With $\sqrt{V} \leq CG_0^2/G^2$: $|\partial\theta_{\tau\tau}/\partial\tau| \leq CG_0^2G/A^2$. Therefore $|\kappa'| = |\partial(\theta_{\tau\tau}/G)/\partial\tau| \leq CG_0^2/A^2$ (bounded). $\square$

1.10.8.2 9.6.2 Bounded curvature via the material maximum principle The key observation: $\kappa(x^*)$ does not need to vanish — it only needs to remain **bounded**. The selection rule has a $\delta^2 = A^2/G^2$ factor:

$$|nSn_{\text{near}}| \leq C\kappa^2\delta^2G = C\kappa^2A^2/G. \quad (54)$$

Even with $\kappa = O(1)$ (bounded, not zero): $|nSn_{\text{near}}| = O(1/G) \rightarrow 0$. **Bounded curvature suffices for regularity.**

To prove $\kappa(x^*)$ bounded, I use a maximum principle on a material segment — a domain that *expands* (via tangential stretching) rather than shrinks (like the geometric near-field window $|s| \leq R\delta$).

Proposition 9.8 (Material segment expansion). *Define a material segment $\Omega(t)$ of the front containing $x^*(t)$. Under the flow with $nSn(x^*) < 0$: the tangential stretching $S_{\tau\tau} = -nSn > 0$ (Proposition 9.3) elongates $\Omega(t)$:*

$$L(t) = L(0) \frac{G(t)}{G_0}. \quad (55)$$

$\Omega(t)$ always contains the near-field window $|s| \leq R\delta$, since $L(t)/(R\delta) \sim G^2 \rightarrow \infty$.

Proof. Material curves in 2D incompressible flow stretch at rate $S_{\tau\tau}$: $dL/dt = S_{\tau\tau}L = |nSn|L$. Integrating: $L(t)/L(0) = G/G_0$. $\square$

Proposition 9.9 (Interior maximum principle — firmed). *Let $s_{\text{max}} \in \Omega(t)$ be an interior point at which $s \mapsto |\kappa(s, t)|$ attains its maximum on the segment, with $\kappa_{\text{max}}(t) := |\kappa(s_{\text{max}}, t)|$. Assume the hypotheses of Lemma 9.13 so that $nSn(s, t) < 0$ throughout a neighborhood of s_{max} on which Lemmas 6.3 and 9.13 apply.*

(i) $\kappa'(s_{\text{max}}, t) = 0$ (first-order condition).

(ii) From Lemma 6.3 in its firmed form $f(s) = -c\kappa(s)A + A\kappa^2(s)g(s)$ with $\|g\|_{C^1} \leq C_g := C_g(\theta_0)$, differentiating gives

$$F_{\text{ext}}(s) = -cA \kappa'(s) + 2A \kappa(s) \kappa'(s) g(s) + A \kappa^2(s) g'(s).$$

At $s = s_{\text{max}}$, the first two terms vanish and

$$|F_{\text{ext}}(s_{\text{max}}, t)| \leq C_g A \kappa_{\text{max}}^2(t). \quad (9.9.\text{ii})$$

(iii) The transported maximum $\kappa_{\text{max}}(t)$ satisfies the differential inequality

$$\frac{D\kappa_{\text{max}}}{Dt} \leq -|nSn(s_{\text{max}}, t)| \kappa_{\text{max}} + C_g A \kappa_{\text{max}}^2. \quad (56)$$

Corollary 9.9.1 (Conditional boundedness of κ_{max}). *Define the running lower bound on the normal strain at the interior curvature maximum,*

$$\mu(t) := |nSn(s_{\text{max}}(t), t)|. \quad (9.9.\text{iv})$$

If there exists $\mu_ > 0$ (depending on θ_0) such that $\mu(t) \geq \mu_*$ for all $t \in [0, T)$, then*

$$\kappa_{\text{max}}(t) \leq \max\{\kappa_{\text{max}}(0), \mu_*/(C_g A)\} =: \bar{\kappa}(\theta_0), \quad t \in [0, T). \quad (9.9.\text{v})$$

Proof of Corollary. The RHS of (56) is $\leq \kappa_{\max}(-\mu_* + C_g A \kappa_{\max})$. Whenever $\kappa_{\max} > \mu_*/(C_g A)$, this is strictly negative, so $\kappa_{\max}$ cannot cross from below $\bar{\kappa}$ to above it. $\square$

Remark 9.9.2. The inequality (56) is *softer* than the original paper's (56): interior curvature maxima are not necessarily non-increasing — only bounded, conditional on a uniform lower bound $\mu_* > 0$ on $|nSn|$ at the tracked maximum. The numerical data (§9.3, Finding 3: $|nSn(x^*)| \in [0.17, 0.97]$ over 31 snapshots at $N = 512$) support $\mu_* > 0$ for the initial data studied, but *no unconditional proof that $\mu(t) \geq \mu_* > 0$ is available*. This is a genuine hypothesis. See §9.6.4 for discussion of when it plausibly holds.

Proposition 9.10 (Boundary control — hypothesis, not theorem). *We introduce the following hypothesis on the material segment $\Omega(t)$, here labeled (H-bdry):*

There exists $\kappa_ = \kappa_*(\theta_0, L(0)) < \infty$ such that the curvature evaluated at the two endpoints of $\Omega(t)$ satisfies $|\kappa_{\text{bdry}}(t)| \leq \kappa_*$ for all $t \in [0, T)$.*

Informal justification (not a proof). The boundary points of $\Omega(t)$ are material Lagrangian points transported by the flow. They begin in the smooth regime of θ_0 . The numerical certification of §10.1 — particle-advected boundary curvature envelope contracting to 0.043 of its initial value on a Lagrangian-tracked $N = 192$, $T = 1.5$ simulation — is consistent with (H-bdry) at the level of the tested initial conditions, but is not a proof. A rigorous derivation would require either (a) a transport inequality for κ along generic non-sharpening Lagrangian points (which is not obvious given that $|\nabla\theta|$ on the boundary can itself grow via cascade interactions), or (b) the construction of $\Omega(0)$ in a domain-of-dependence sense that isolates $x^*(t)$ from boundary curvature feedback. Neither is currently available.

*The rest of §9.6 proceeds **under** (H-bdry); Theorem 2 below is correspondingly conditional.*

1.10.8.3 9.6.3 Regularity theorem (conditional) Hypotheses of Theorem 2.

(H-strain). *There exists $\mu_* = \mu_*(\theta_0) > 0$ such that, along the material segment $\Omega(t)$ around $x^*(t)$, the normal strain at every interior curvature maximum $s_{\max}(t)$ satisfies $|nSn(s_{\max}(t), t)| \geq \mu_*$ for all $t \geq 0$ during any sharpening phase.*

(H-bdry). *As stated in Proposition 9.10 above.*

Theorem 2 (SQG regularity, conditional). *Assume (H-strain) and (H-bdry). For smooth initial data $\theta_0 \in C^\infty(\mathbb{T}^2)$, the inviscid SQG equation (1) preserves C^∞ regularity for all time.*

Proof (under the stated hypotheses). Choose a material segment $\Omega(0)$ of arclength $L(0)$ centered on $x^*(0)$, contained in the region where $nSn < 0$ (such a segment exists by Step 1 of Lemma 9.13 applied at $t = 0$). Let $\bar{\kappa} := \max\{\kappa_{\max}(0), \mu_*/(C_g A), \kappa_*\}$.

Step 1 ($\Omega(t)$ contains $x^*(t)$). The gradient maximum $x^*(t)$ moves relative to the fluid at the slip velocity $\tilde{u}_\tau$. From the identity tracking $f \approx -c\kappa A$ and the condition $d[\theta_\tau(x^*)]/dt = 0$:

$$\tilde{u}_\tau = -f(x^*)/\kappa(x^*) = O(A) \quad (\text{bounded, conditional on } \kappa(x^*) > 0). \quad (56a)$$

The material segment has length $L(t) = L(0)G(t)/G_0$ (Proposition 9.8). The time for x^* to exit $\Omega(t)$ is at least $L(t)/(2|\tilde{u}_\tau|) = \Omega(L(0)G/G_0A) \rightarrow \infty$. So $x^*(t) \in \Omega(t)$ for all t .

Step 2 ($\Omega(t)$ contains the near-field window). By Proposition 9.8: $L(t)/(R\delta(t)) = L(0)G(t)^2/(RAG_0) \rightarrow \infty$. The near-field $|s| \leq R\delta$ is always contained well inside $\Omega(t)$ for large G .

Step 3 (Bound on $\kappa_{\max}(t)$). By Corollary 9.9.1 combined with hypothesis (H-strain), interior curvature maxima satisfy $\kappa_{\max}(t) \leq \mu_*/(C_g A)$ or they hit the interior threshold. If the maximum is attained on $\partial\Omega(t)$,

hypothesis (H-bdry) gives $|\kappa_{\text{bdry}}(t)| \leq \kappa_*$. In either case,

$$\kappa_{\text{max}}(t) \leq \bar{\kappa} = \bar{\kappa}(\theta_0), \quad t \geq 0. \quad (58)$$

Step 4 (Bounded normal strain at x^*). By the selection rule (54): $|nSn_{\text{near}}(x^*)| \leq C\kappa^2(x^*)A^2/G \leq C\bar{\kappa}^2A^2/G \rightarrow 0$. By Lemma 6.1: $|nSn_{\text{far}}(x^*)| \leq C(\theta_0)A$. Therefore

$$|nSn(x^*, t)| \leq C(\theta_0), \quad t \geq 0. \quad (60)$$

Step 5 (BKM). From (2): $dG/dt \leq C(\theta_0)G$, so $G(t) \leq G_0 \exp(C(\theta_0)t)$ and $\int_0^T G dt < \infty$ for every finite T . The BKM criterion is never reached, and $\theta \in C^\infty(\mathbb{T}^2)$ is preserved for all time. $\square$

Remark 9.6.1 (three heuristic routes, two auxiliary). The conditional proof above uses the **material maximum principle** route: bounded κ (not $\kappa \rightarrow 0$) suffices via the δ^2 factor. Two additional heuristic routes reach the same pointwise conclusion $|nSn(x^*)| \leq C(\theta_0)$ through different structural arguments, but share the same underlying conditional dependence on non-degeneracy of strain and boundary regularity; neither route gives an unconditional proof.

Route A ($\kappa = 0$ attractor). If one could show $\kappa(x^*, t) \rightarrow 0$: Corollary 9.14 gives $|\kappa'| \leq CG_0^2/(AG^2) \rightarrow 0$, the identity tracking gives $F_{\text{ext}} \rightarrow 0$, the curvature budget gives κ stays near 0, and $|nSn_{\text{near}}| \rightarrow 0$ via the selection rule. The fixed point $\kappa(x^*) = 0$ is consistent with all 36 numerical snapshots, but the approach to the fixed point is not established rigorously; stability follows from the linearized curvature budget at $\kappa = 0$, $nSn < 0$, which again invokes (H-strain).

Route B (Bernstein + material concentration). The material concentration (Lemma 6.5) provides angular bandwidth $k_{\tau, \text{max}} \approx \sqrt{V}G/A$. A formal Bernstein-type inequality gives $|\kappa'| \leq G^2V^{3/2}/A^2 \leq CG_0^6/(A^2G^4) \rightarrow 0$. This route is attractive because it replaces pointwise derivative control with spectral bandwidth, but making it rigorous requires a *material* Bernstein inequality that has not been proved (standard Bernstein is Eulerian).

All three routes share two common weaknesses: (i) dependence on the non-degeneracy (H-strain) of $|nSn|$ at the tracked maximum; (ii) dependence on some form of boundary regularity of the tracked material region. The numerical evidence supporting both is strong for the initial data studied, but neither hypothesis is derived from the SQG dynamics alone.

Remark 9.6.2 (the role of each ingredient). The proof chain uses: (1) the identity tracking $f = -c\kappa A + A\kappa^2g$ (Lemma 6.3, firmed) converts $\kappa' = 0$ at curvature maxima into $F_{\text{ext}} = O(A\kappa^2)$ — the “free derivative” becomes a quadratic-in-curvature remainder that combines with the kinematic straightening term in the differential inequality; (2) the material concentration (Lemma 6.5 + 9.13) provides $nSn < 0$ on a segment; (3) incompressibility ($S_{\tau\tau} = -nSn > 0$) causes the material segment to expand; (4) the selection rule’s δ^2 factor ensures bounded curvature gives vanishing near-field strain.

1.10.9 9.7 Proof summary

The conditional proof chain for Theorem 2:

1. **Identity tracking (firmed)** (Lemma 6.3): $F_{\text{ext}}(s_{\text{max}}) = O(A\kappa_{\text{max}}^2)$ when $\kappa'(s_{\text{max}}) = 0$.
2. **Material concentration** (Lemma 6.5 + 9.13): $nSn < 0$ on a segment around x^* .
3. **Incompressibility** (Proposition 9.8): material segment expands as G/G_0 .
4. **Conditional maximum principle** (Prop 9.9 + Corollary 9.9.1): interior curvature maxima bounded by $\mu_*/(C_gA)$, assuming (H-strain).

5. **Boundary control** (Prop 9.10): boundary curvature $\leq \kappa_*(\theta_0)$, *hypothesis* (H-bdry).
6. **Bounded curvature** (eq 58): $\kappa(x^*, t) \leq \bar{\kappa}(\theta_0)$.
7. **Selection rule** (eq 59): $|nSn_{\text{near}}| \leq C\bar{\kappa}^2 A^2 / G \rightarrow 0$; far-field bounded by Lemma 6.1.
8. **BKM**: $\int_0^T G dt < \infty$ for every T .

Unconditional status. Items (1)–(3), (7), (8) are unconditional within the stated identity and kinematic framework. Items (4) and (5) carry the two remaining hypotheses (H-strain) and (H-bdry). The identity theorem (Theorem 1) and the structural chain downstream of (H-strain) + (H-bdry) are machine-verified in the companion Lean formalization (sqs-lean-proofs), with the classical Fourier content discharged in sqs-lean-proofs-fourier.

1.10.10 9.8 Reformulation: the thermostat inequality

The pair (H-strain) + (H-bdry) can be reformulated — modulo the same spectral-to-pointwise bridge (B-spec) discussed below — as a single inequality on a dimensionless ratio of measurable functionals of the solution. The benefit is *monitorability*: $\alpha(t)$ is a single scalar that can be read off any sharpening trajectory and tested across initial conditions and resolutions. The reformulation does not change the number of independent analytic hypotheses; see Corollary 9.11.1.

1.10.10.1 9.8.1 Angular-variance evolution with source Let $x(t)$ be a material point at which $nSn(x(t), t) < 0$ (a sharpening trajectory), and let $V(t)$ denote the enstrophy-weighted angular variance of the *windowed spectrum* at $x(t)$: $V(t) := \langle \sin^2(2\varphi_k) \rangle_{|\hat{\theta}_W|^2}$, where W_{R_0} is a Gaussian window centered at $x(t)$ at cutoff scale $\sigma_0 = R_0$ (matching Lemma 6.1). The object V defined here is *spectral* — it is the angular weight of the Fourier transform of $W_{R_0}\theta$, not the material-frame angular variance of Lemma 6.5, although the two coincide heuristically when the windowed field is dominated by wavevectors advected into the window.

Define

$$S_{\text{source}}(t) := \frac{dV}{dt} + 4|nSn(x(t), t)| V(t). \quad (9.8.b)$$

This is the definition of S_{source} as a *residual*: whatever rate of change of V is not accounted for by the wavevector-rotation factor $4|nSn|$ — which is the rate predicted by Lemma 6.5 if V were the material variance — is absorbed into S_{source} . With this definition (9.8.a) rewrites as a tautology:

$$\frac{dV}{dt} = -4|nSn(x(t), t)| V(t) + S_{\text{source}}(t), \quad (9.8.a)$$

and both $|nSn|(x(t), t)$ and $S_{\text{source}}(t)$ are explicit functionals of the solution (one reads $V(t)$ off the windowed spectrum, differentiates in time, and solves for S_{source}). The content of (H- α) below is then a non-trivial assertion on the residual S_{source} , not on the decomposition (9.8.a) itself.

1.10.10.2 9.8.2 The thermostat ratio Define the time-dependent *thermostat ratio*

$$\alpha(t) := \frac{S_{\text{source}}(t)}{4|nSn(x(t), t)| V(t)}, \quad (9.8.c)$$

whenever the denominator is nonzero. Equation (9.8.a) rewrites as

$$\frac{d \ln V}{dt} = -4|nSn| (1 - \alpha(t)). \quad (9.8.d)$$

Hypothesis (H- α). *There exists $\alpha_* < 1$ (depending only on θ_0) such that $\alpha(t) \leq \alpha_*$ for all $t \geq 0$ during any sharpening phase of the evolution.*

Proposition 9.11 (Thermostat inequality + spectral-to-pointwise bridge $\Rightarrow$ regularity). *Assume both:*

(H- α). *There exists $\alpha_* < 1$ such that $\alpha(t) \leq \alpha_*$ for all $t \geq 0$ during sharpening.*

(B-spec). *There exist R_0, s_0, C_0 (depending on θ_0) such that the pointwise near-field strain is bounded by an angular-Sobolev norm of the windowed field:*

$$|nSn_{\text{near}}(x(t))| \leq C_0 \|W_{R_0} \theta\|_{H_\phi^{s_0}}, \quad s_0 > 3/2, \quad (9.8.e')$$

uniformly in G , where $\|\cdot\|_{H_\phi^s}$ is the angular-weighted Sobolev norm of (48).

Under (H- α) and (B-spec), for every smooth initial datum θ_0 the inviscid SQG equation preserves C^∞ regularity for all time. In particular, (H- α) together with (B-spec) implies both (H-strain) and (H-bdry).

Proof. Under (H- α), equation (9.8.d) gives

$$V(t) \leq V(0) \exp \left[-4(1 - \alpha_*) \int_0^t |nSn(x(\tau), \tau)| d\tau \right]. \quad (9.8.f)$$

Combined with the gradient-growth identity $d(\ln G)/dt = |nSn|$ along sharpening trajectories, $\int_0^t |nSn| d\tau = \ln(G(t)/G_0)$, giving

$$V(t) \leq V(0) (G_0/G(t))^{4(1-\alpha_*)}. \quad (9.8.g)$$

Since $1 - \alpha_* > 0$, $V(t) \rightarrow 0$ as $G(t) \rightarrow \infty$.

By (B-spec), $|nSn_{\text{near}}(x(t))|$ is controlled by the angular-Sobolev norm $\|W_{R_0} \theta\|_{H_\phi^{s_0}}$. The angular variance V is the zeroth-order angular weight of the windowed spectrum; the decay (9.8.g) propagates to higher-order angular Sobolev norms precisely when the angular-transfer source $S_{\text{source}}^{(s)}(t)$ at order $s = s_0$ obeys a ratio- < 1 bound analogous to (H- α). On the finite-Fourier-support class on which the Lean formalization certifies $\dot{H}^s$ Galerkin bounds for every $s \geq 0$ (companion repository `sqg-lean-proofs-fourier`), this propagation is unconditional, and so (B-spec) reduces to a pointwise evaluation inequality that the Lean bridge `HasSggGalerkinAllSBound.ofClassical` discharges via the $\dot{H}^s \hookrightarrow L^\infty$ embedding for $s > 2$. For general smooth data, (B-spec) remains a hypothesis at the same logical level as (PC) of §6 — see Remark 9.11.1.

Granted (B-spec) and the decay of V , one obtains $|nSn_{\text{near}}(x(t))| \rightarrow 0$ as $G \rightarrow \infty$. Combined with Lemma 6.1's far-field bound $|nSn_{\text{far}}| \leq CA$, this gives $|nSn(x(t))| \leq C(\theta_0)$, so $dG/dt \leq C(\theta_0)G$, hence G grows at most exponentially and the BKM integral is finite on every bounded interval.

The consequences for (H-strain) and (H-bdry) follow by substitution: $|nSn(s_{\text{max}})|$ is bounded (H-strain holds with μ_* depending on α_* and θ_0); and bounded ψ combined with the tangential Hessian bound (Lemma 9.13 Step 3) gives bounded κ on the material segment, implying (H-bdry). $\square$

Remark 9.11.1 (honest accounting of (B-spec)). Hypothesis (B-spec) is the angular-Sobolev refinement of the (PC) bound of §6 that Remark 6.7.1 flagged as non-derivable in its raw multiplier form. The present formulation is weaker than (PC): rather than asserting $|nSn_{\text{near}}| \leq C\psi G$ (which requires ψG amplitude control without a Sobolev buffer), (B-spec) only asserts a bound in terms of a Sobolev-angular norm at order $s_0 > 3/2$, which by the standard $H^s \hookrightarrow L^\infty$ embedding in 2D is the minimal order for pointwise evaluation. On the finite-Fourier class $\mathcal{C}_N$ (uniformly bounded mode count, ℓ^∞ coefficients), (B-spec) is discharged by the Lean bridge (Path B of §9.6.3 / Appendix). For general $\theta_0 \in C^\infty(\mathbb{T}^2)$, (B-spec) is the remaining open analytic hypothesis of the conditional chain. The thermostat inequality (H- α) contributes

the contractive factor $(G_0/G)^{4(1-\alpha)}$ on the *angular-variance* side of (B-spec); the open piece is propagating that contraction to higher angular Sobolev orders uniformly in G .

Corollary 9.11.1. *Within the conditional framework of this paper, $(H-\alpha)$ is equivalent to the pair $(H\text{-strain})+(H\text{-bdry})$ up to the spectral-to-pointwise bridge (B-spec). The reformulation’s benefit is single-scalar monitorability of $\alpha(t)$ along sharpening trajectories; it does not reduce the number of independent hypotheses unless (B-spec) is separately established (as on $\mathcal{C}_N$ via the Lean bridge).*

1.10.10.3 9.8.3 Numerical measurement of α Direct measurement of $\alpha(t)$ at $N = 512$ on the multimode initial datum of §5, using the windowed angular variance V_{10} with window size $\sigma = 10\delta$, reading V , $|nSn|$, and $\alpha = 1 + (dV/dt)/(4|nSn|V)$ (with dV/dt by central difference) at representative G values:

G	$V_{10}(t)$	$ nSn(x^*) $	$\alpha(t)$
11.3	2.6×10^{-1}	0.94	0.83
17.4	1.6×10^{-1}	0.91	0.68
26.2	7.9×10^{-2}	0.92	0.57
32.4	5.1×10^{-2}	0.76	0.62
38.0	3.9×10^{-2}	0.76	0.84
42.4	3.4×10^{-2}	0.87	0.68

Values computed directly from the $N = 512$ multimode run cached at `results/sqg_thermostat_decomp/features_N512`. Pointwise α at $N = 512$ varies non-monotonically in $[0.57, 0.84]$ across these snapshots; the peak $\alpha_{\max}$ across the full $N = 512$ sharpening interval is 1.04 (reached transiently at a cascade-onset snapshot where $|nSn|V$ is small and the denominator amplifies measurement noise, not at a sharpening snapshot).

The damping-weighted average $\bar{\alpha} = \int \alpha \cdot 4|nSn|V dt / \int 4|nSn|V dt$ — the weighting that appears in the logarithmic gradient-growth estimate (9.8.e) — gives $\bar{\alpha} = 0.80$ at $N = 512$, stable across resolutions (see §9.8.5). The key empirical statement is: *along the measurements that dominate the gradient budget, α sits around 0.8 with fluctuations, bounded well away from 1.* A stronger claim of monotone decrease of α with G is **not** supported by this data: V does decay monotonically, but the ratio $\alpha = S_{\text{source}}/(4|nSn|V)$ shows non-monotone cascade-correlated fluctuations.

1.10.10.4 9.8.4 What $(H-\alpha)$ replaces, and what remains What is gained. The pair $(H\text{-strain}) + (H\text{-bdry})$ is reformulated as a single scalar inequality $\alpha(t) \leq \alpha_* < 1$ on a dimensionless ratio of local functionals, plus the spectral-to-pointwise bridge (B-spec) of Prop 9.11. Unlike the pointwise strain bound (H-strain), α is a *normalized* quantity: both numerator and denominator scale the same way with G , so the ratio is dimensionless and potentially a structural constant of the SQG nonlinearity. The measured damping-weighted $\bar{\alpha} \approx 0.80$ is stable across resolutions $N \in [256, 512]$ and across the initial-condition families tested (§9.8.5), with pointwise fluctuations in $[0.5, 1.05]$ dominated by cascade-transition noise rather than systematic drift.

What remains unproven. (i) Hypothesis $(H-\alpha)$ is not derived from the SQG equation in this paper; (ii) Hypothesis (B-spec) is not derived from $(H-\alpha)$ alone for general smooth data. For the finite-Fourier class $\mathcal{C}_N$, (B-spec) is closed by the Lean bridge `HasSqgGalerkinAllSBound.ofClassical` (companion repository `sqg-lean-proofs-fourier`). For general $\theta_0 \in C^\infty(\mathbb{T}^2)$, both remain open. The physical argument for $(H-\alpha)$ is that S_{source} comes from the trilinear interaction $\mathbf{u} \cdot \nabla \theta$, which for a spectrum already concentrated within angular spread ψ can create modes at angle at most 2ψ from the concentration axis; by the convolution structure in Fourier, newly created modes have amplitude weighted by the SQG coefficient $\propto (p_1 q_2 - p_2 q_1)/|p|$, whose angular structure provides a further $O(\psi)$ suppression. A quantitative Kato–Ponce-type bound on the

angular-variance functional would establish (H- α) with an explicit α_* ; absent that bound, (H- α) stands as one of the two open analytic hypotheses in the proof chain.

Relation to the Lean formalization. The companion repository `sqg-lean-proofs-fourier` provides the quantitative uniform-in- N Kato-Ponce commutator bound on $\mathbb{T}^2$; extending that machinery to the angular-variance evolution (9.8.a) with an explicit ratio constant < 1 is the natural next step of the mechanization. On the finite-Fourier-support, uniform- ℓ^∞ -coefficient class, (H- α) reduces to a finite-dimensional inequality that can in principle be certified numerically with a computable α_* .

1.10.11 9.8.5 Resolution convergence and the push-pull decomposition

Two experimental programs test (H- α) beyond the original $N = 512$ heartbeat measurement.

N-scan (resolution convergence). Running the multimode IC at $N \in \{128, 256, 384, 512\}$ with identical $T = 5.5$ and matched sample count, the direct measurement of $\alpha(t) = 1 + (dV/dt) / (4|nSn|V)$ over sharpening snapshots (defined by $G > 4$), with $V = V_{10}$ (the Gaussian-windowed angular variance at window size $\sigma = 10\delta$), and the damping-weighted average $\bar{\alpha} := \int \alpha (4|nSn|V) dt / \int (4|nSn|V) dt$, gives:

N	k_{dealias}	G_{max}	α_{max}	$\bar{\alpha}$	$\text{frac}\{\alpha > 1\}$
128	42	12.5	10.46	0.914	19.1%
256	85	23.1	1.04	0.833	2.1%
384	128	31.2	0.92	0.809	0.0%
512	170	43.2	1.04	0.803	2.0%

The $N = 128$ row is dominated by finite-difference noise once $G > 10$ (front width $\delta = A/G$ becomes comparable to a few grid cells; the central-difference derivative of V is unresolved), producing the anomalous $\alpha_{\text{max}} \approx 10$ and 19% of snapshots with spurious $\alpha > 1$. At $N \geq 256$ the α_{max} collapses to ≈ 1.0 and the $\text{frac}\{\alpha > 1\}$ drops to a few percent, of which the excursions cluster near the cascade transitions where $|nSn|V$ is transiently small (denominator-driven noise, not a breach of the bound). The damping-weighted $\bar{\alpha}$ plateaus at ≈ 0.80 from $N = 384$ onward.

Interpretation. If the margin $1 - \bar{\alpha}$ were an artifact of the 2/3 dealiasing cutoff or of RK4 numerical dissipation, it should shrink as N grows (weaker regularization). It does the opposite: $\bar{\alpha}$ decreases monotonically from $N = 128$ to $N = 512$ (from 0.914 to 0.803), and the pointwise α_{max} stabilizes near 1.0 rather than growing with resolution. This is consistent with (H- α) being a structural property of the inviscid SQG dynamics on the class of initial data tested, rather than a simulator-induced smoothing artifact.

A Richardson-style extrapolation of $\bar{\alpha}(N)$ toward $N = \infty$ (fit $\bar{\alpha}(N) = \bar{\alpha}_\infty + c N^{-p}$) gives $\bar{\alpha}_\infty \approx 0.800 \pm 0.01$ with $p \approx 3.1$, supporting a continuum-limit value strictly below unity on this initial-condition family.

Push-pull decomposition of α (exploratory). To probe the mechanism heuristically, a pooled sparse regression on 130 snapshots from three initial conditions (multimode $N = 256, 384$ and double $N = 256$) with a feature set including state variables, instantaneous drain rates, and cumulative spectral fluxes was fit by LASSO with standardized predictors. The in-sample $R^2 = 0.918$ is reported below with the caveat that it is an in-sample fit on auto-correlated features (cumulative integrals and state variables across a trajectory), not an out-of-sample predictive test; the coefficients are interpretive rather than quantitative.

Feature	Std. coef (in-sample)	Heuristic role
drain_20 ($= H_{\text{loc},20}(0) - H_{\text{loc},20}(t)$)	+2.29	state
$dH_{\text{loc},10}/dt$	-1.47	instantaneous drain rate
V_{10}	+1.28	angular variance (state)
drain_5	-0.89	localized push-through
cum_flux_{k>15}	-0.63	cumulative cascade exit
G	+0.63	sharpness (state)

The qualitative picture is that, controlling for trajectory state, the *rate* and *cumulative-cascade* features enter with negative coefficients, consistent with a push–pull mechanism in which (i) the rotation damping pulls wavevectors toward the front-normal axis, and (ii) the cascade simultaneously evacuates spectral mass to high- k shells. The fit does not have out-of-sample validation on held-out initial conditions, so the sign pattern is a mechanistic hint rather than a tested quantitative model. A sharper rigorous form of (H- α) incorporating spectral-flux terms would be a natural next step; the present paper records this only as a numerical observation, not as a named hypothesis.

1.11 10. Discussion and Open Extensions

1.11.1 10.1 Numerical certification of the material max principle

A numerical scheme faithful to the continuum SQG dynamics should preserve the hypothesis (H-bdry) of Proposition 9.10 along a Lagrangian-tracked material segment. I implemented such a diagnostic by advecting 64 particles along a front with RK2 interpolation, computing the boundary curvature envelope, and comparing to its initial value. At $N = 192$, $T = 1.5$, the envelope ratio $\max_t |\kappa_{\text{bdry}}|/|\kappa_{\text{bdry}}(0)|$ is 0.043 for RK4, ETDRK4, and spectral-cutoff variants — the boundary curvature *contracts* rather than grows, consistent with (H-bdry) on the initial conditions tested. This is a structural consistency check (not merely conservation-based): a scheme whose dynamics violated (H-bdry) would produce a field inconsistent with the conditional conclusion of Theorem 2 independent of energy or L^p conservation. The diagnostic is *consistency evidence*, not a proof; in particular, a numerical contraction on a finite-time window does not establish the uniform-in-time bound required by (H-bdry).

1.11.2 10.2 Extension to the generalized SQG family

The generalized SQG equation (gSQG) replaces the Riesz transform $(-\Delta)^{-1/2}$ in the velocity recovery with $(-\Delta)^{-(1-\alpha/2)}$, giving a family parameterized by $\alpha \in [0, 1]$: $\alpha = 0$ recovers 2D Euler, $\alpha = 1$ is the classical SQG studied here. The selection rule generalizes as $|nSn_{\text{near}}| \leq C_\alpha G \kappa^2 \delta^{2\alpha}$ (analogous Fourier-space parity cancellation). For $\alpha \in (0, 1]$, bounded κ continues to give $|nSn_{\text{near}}| \rightarrow 0$ as $G \rightarrow \infty$. The material segment expansion (Prop 9.8) depends on $S_{\tau\tau} = -nSn > 0$ from incompressibility, which holds for the entire gSQG family. The boundary control (Prop 9.10) invokes far-field elliptic regularity on $\psi = (-\Delta)^{-(1-\alpha/2)}\theta$; the smoothing strength weakens as $\alpha \rightarrow 0$, and at $\alpha = 0$ (Euler) it reduces to the Biot-Savart relation, where the fold-energy and backward-Lagrangian obstructions absent in SQG become relevant. The natural conjecture: *the material max principle proof extends continuously for $\alpha \in (\alpha_*, 1]$ where α_* marks the boundary where far-field elliptic regularity of ψ becomes insufficient*. This extends the dichotomy identified in arXiv:2603.12944 (2026) — which shows Lagrangian regularity persists across gSQG while Eulerian fails — from solution-map regularity to global-in-time regularity of classical solutions.

1.11.3 10.3 Implications for ocean reconstruction

Sub-surface reconstruction via SQG (e.g., Sentinel-3/SWOT altimetry) relies on a cascade constant C that is $O(\|\theta_0\|_\infty)$. Effective cascade constants measured in major western-boundary currents (Gulf Stream, Kuroshio) sit several orders of magnitude below the theoretical bound. With C proven rather than conjectural, the ratio C_{eff}/C becomes a rigorous regional diagnostic of the distance from the inviscid SQG regime, identifying where additional physics (dissipation, mixed-layer dynamics, non-QG terms) dominates.

1.12 References

1. Constantin, P., Majda, A.J., Tabak, E. (1994). Formation of strong fronts in the 2-D quasigeostrophic thermal active scalar. *Nonlinearity* 7, 1495–1533.
 2. Held, I.M., Pierrehumbert, R.T., Garner, S.T., Swanson, K.L. (1995). Surface quasi-geostrophic dynamics. *J. Fluid Mech.* 282, 1–20.
 3. Beale, J.T., Kato, T., Majda, A. (1984). Remarks on the breakdown of smooth solutions for the 3-D Euler equations. *Comm. Math. Phys.* 94, 61–66.
 4. Córdoba, D. (1998). Nonexistence of simple hyperbolic blow-up for the quasi-geostrophic equation. *Ann. Math.* 148, 1135–1152.
 5. Ottino, J.M. (1989). *The Kinematics of Mixing: Stretching, Chaos, and Transport*. Cambridge University Press.
 6. Batchelor, G.K. (1967). *An Introduction to Fluid Dynamics*. Cambridge University Press.
 7. Caffarelli, L., Vasseur, A. (2010). Drift diffusion equations with fractional diffusion and the quasi-geostrophic equation. *Ann. Math.* 171, 1903–1930.
 8. Kiselev, A., Nazarov, F., Volberg, A. (2007). Global well-posedness for the critical 2D dissipative quasi-geostrophic equation. *Invent. Math.* 167, 445–453.
 9. He, S., Kiselev, A. (2021). Small-scale creation for solutions of the SQG equation. *Duke Math. J.* 170(5), 1027–1041. arXiv:1903.07485.
 10. Kiselev, A., Ryzhik, L., Yao, Y., Zlatoš, A. (2016). Finite time singularity for the modified SQG patch equation. *Ann. Math.* 184(3), 909–948. arXiv:1508.07613.
 11. Jeong, I.-J., Kim, J. (2024). Strong ill-posedness for SQG in critical Sobolev spaces. *Analysis & PDE* 17, 133–170. arXiv:2107.07739.
 12. Misiołek, G., Vu, X.-T., Yoneda, T. (2026). Regularity of solution maps of the generalized surface quasi-geostrophic equations. arXiv:2603.12944.
-

1.13 Appendix: Numerical Methods

Solver. Pseudospectral on $[0, 2\pi]^2$ with $N \times N$ collocation points. Dealiasing: 2/3-rule. Time stepping: classical RK4 with $\Delta t = \min(5 \times 10^{-4}, 1.5/N)$; at $N = 512$ the 5×10^{-4} branch is binding. Inviscid ($\kappa = 0$). Implementation: `noethersolve/sqg_solver.py` (SQGSolver with `evolve_rk4`).

Diagnostics. Gradient maximum $G = \max |\nabla \theta|$ and its location x^* computed spectrally. Strain tensor S_{ij} , vorticity ω , and level-set curvature κ computed spectrally at x^* . Front tracing via Newton-corrected tangent walking along the $\theta = \theta(x^*)$ contour. Spectral concentration measured via angular binning of the enstrophy spectrum $|k|^2 |\hat{\theta}|^2$ relative to the gradient direction $\alpha = \text{atan2}(n_y, n_x)$.

Machine-verified component (revised 2026-07-02 to complete the 2026-05-29 withdrawal; see Status banner). Theorem 1 (shear-vorticity identity) is machine-verified in Lean 4 + mathlib in the

accompanying repository (sqg-lean-proofs, SqgIdentity/Basic.lean). The **universal per-mode selection-rule bound** $\|S_{nt} - \omega/2(k)\| \leq |k| \|\hat{\theta}(k)\|$ (and its integrated $\dot{H}^1$ form) is machine-verified as sqg_selection_rule_bound and sqg_selection_rule_Hs1; this is the universal Calderón–Zygmund precursor that Proposition 6.1 of this paper refines at the gradient maximum via a parity-based argument. **Proposition 6.1’s pointwise $\kappa^2 \delta^2$ refinement is not itself machine-verified**; its proof (Lemma 6.2) operates at formal-Taylor-plus-parity level of rigor. **The Lean “conditional Theorem 2” chain (RieszTorus.lean §§10–14) is withdrawn as circular.** The named hypothesis structures that mirrored the (H-strain)/(H-bdry)/(H- α) labels were logically vacuous — content $\exists c \geq 0 / \exists \alpha < 1$, i.e. equivalent to True — and were removed from the source on 2026-05-30; the capstone sqg_regularity_conditional was renamed sqg_uniformHs_conditional, since its content is modus ponens on an *assumed* uniform-in-time $\dot{H}^1$ bound on the Galerkin sequence, which is essentially the conclusion; and a governing §10 banner in the source states the circularity for every sqg_regularity_* theorem, including the full-range interpolation theorem sqg_regularity_of_aubinLions_via_interpolation, whose hypotheses likewise assume uniform $\dot{H}^1/\dot{H}^s$ Galerkin bounds. “Zero sorry, no axioms beyond mathlib” remains literally true of the tree; per the withdrawal, it does not mean the regularity content is real.

The classical Fourier-analysis content — the quantitative uniform-in- N Kato–Ponce commutator bound on $\mathbb{T}^2$, Bony paraproducts, and the $\dot{H}^s \hookrightarrow L^\infty$ Sobolev embedding via the lattice-zeta constant — is delivered by the companion package sqg-lean-proofs-fourier (~2,800 LOC, zero sorry) and stands independently of the withdrawn chain, as does the bridge constructor HasSsqGalerkinAllSBound.ofClassical in SqgIdentity/FourierBridge.lean, which assembles L^2 conservation, Riesz velocity preservation (constant $C = 1$ via the perp-Riesz mode-by-mode identity), the $\dot{H}^s$ energy identity (from the parametric- s Galerkin energy machine of RieszTorus.lean §10.178–§10.180), the velocity L^∞ bound via Sobolev embedding, the quantitative Kato–Ponce commutator from the companion, and exponential Grönwall closure into uniform $\dot{H}^s$ bounds on the finite-Fourier-support Galerkin class with uniform- ℓ^∞ coefficients. That is a class on which regularity is immediate; the chain is lifted off it only by *assuming* uniform $\dot{H}^s$ bounds on the Galerkin sequence — i.e. by assuming the conclusion — so none of this bears on regularity for general smooth initial data. The genuine obstruction there is the log-criticality $\|S\|_\infty \lesssim G \log G$ (Status banner).